

Chapter 4: The SQL Language

**Database Systems
(CO2013)**

Computer Science Program

Assoc. Prof. Dr. Võ Thị Ngọc Châu

[\(chauvtn@hcmut.edu.vn\)](mailto:(chauvtn@hcmut.edu.vn))

Content

- Chapter 1: An Overview of Database Systems
- Chapter 2: The Entity-Relationship Model
- Chapter 3: The Relational Data Model
- **Chapter 4: The SQL Language**
- Chapter 5: Relational Database Design
- Chapter 6: Physical Storage and Data Management
- Chapter 7: Database Security

Chapter 4: The SQL Language

- 4.1. Introduction to the SQL language
- 4.2. DDL
- 4.3. DML
- 4.4. DCL
- 4.5. Stored Functions
- 4.6. Stored Procedures
- 4.7. Triggers

Main References

Text:

- [1] R. Elmasri, S. R. Navathe, Fundamentals of Database Systems- 6th Edition, Pearson- Addison Wesley, 2011.
 - **R. Elmasri, S. R. Navathe, *Fundamentals of Database Systems- 7th Edition, Pearson, 2016.***

References:

- [1] S. Chittayasothorn, *Relational Database Systems: Language, Conceptual Modeling and Design for Engineers*, Nutcha Printing Co. Ltd, 2017.
- [3] A. Silberschatz, H. F. Korth, S. Sudarshan, Database System Concepts – 7th Edition, McGraw-Hill, 2020.
- [4] H. G. Molina, J. D. Ullman, J. Widom, *Database Systems: The Complete Book - 2nd Edition*, Prentice-Hall, 2009.
- [5] R. Ramakrishnan, J. Gehrke, *Database Management Systems – 4th Edition*, McGraw-Hill, 2018.
- [6] M. P. Papazoglou, S. Spaccapietra, Z. Tari, *Advances in Object-Oriented Data Modeling*, MIT Press, 2000.
- [7]. G. Simision, *Data Modeling: Theory and Practice*, Technics Publications, LLC, 2007.

Introduction to the SQL language

| SQL (Structured Query Language) | |
|---------------------------------|--|
| Paradigm | Declarative |
| Family | Query language |
| Designed by | Donald D. Chamberlin Raymond F. Boyce |
| Developer | ISO/IEC |
| First appeared | 1974; 46 years ago |
| Stable release | SQL:2016 / December 2016; 3 years ago |
| Typing discipline | Static, strong |
| OS | Cross-platform |
| Website | www.iso.org/standard/63555.html |

| Major implementations |
|--|
| Many |
| Dialects |
| SQL-86 • SQL-89 • SQL-92 • SQL:1999 • SQL:2003 • SQL:2006 • SQL:2008 • SQL:2011 • SQL:2016 |
| Influenced by |
| Datalog |
| Influenced |
| CQL, LINQ, SPARQL, SOQL, PowerShell, ^[1] JPQL, jOOQ, N1QL |

| SQL (file format) | |
|---------------------|-----------------------------------|
| Filename extension | .sql |
| Internet media type | application/sql ^{[2][3]} |
| Developed by | ISO/IEC |
| Initial release | 1986 |
| Type of format | Database |
| Standard | ISO/IEC 9075 |
| Open format? | Yes |

| Year | Name | Alias | Comments |
|------|----------|------------------|---|
| 1986 | SQL-86 | SQL-87 | First formalized by ANSI. |
| 1989 | SQL-89 | FIPS 127-1 | Minor revision that added integrity constraints, adopted as FIPS 127-1. |
| 1992 | SQL-92 | SQL2, FIPS 127-2 | Major revision (ISO 9075), <i>Entry Level</i> SQL-92 adopted as FIPS 127-2. |
| 1999 | SQL:1999 | SQL3 | Added regular expression matching, recursive queries (e.g. transitive closure), triggers , support for procedural and control-of-flow statements, non-scalar types (arrays), and some object-oriented features (e.g. structured types). Support for embedding SQL in Java (SQL/OLB) and vice versa (SQL/JRT). |
| 2003 | SQL:2003 | | Introduced XML -related features (SQL/XML), window functions , standardized sequences, and columns with auto-generated values (including identity-columns). |
| 2006 | SQL:2006 | | ISO/IEC 9075-14:2006 defines ways that SQL can be used with XML. It defines ways of importing and storing XML data in an SQL database, manipulating it within the database, and publishing both XML and conventional SQL-data in XML form. In addition, it lets applications integrate queries into their SQL code with XQuery , the XML Query Language published by the World Wide Web Consortium (W3C), to concurrently access ordinary SQL-data and XML documents. ^[34] |
| 2008 | SQL:2008 | | Legalizes ORDER BY outside cursor definitions. Adds INSTEAD OF triggers, TRUNCATE statement, ^[35] FETCH clause. |
| 2011 | SQL:2011 | | Adds temporal data (PERIOD FOR) ^[36] (more information at: Temporal database#History). Enhancements for window functions and FETCH clause. ^[37] |
| 2016 | SQL:2016 | | Adds row pattern matching, polymorphic table functions, JSON. |
| 2019 | SQL:2019 | | Adds Part 15, multidimensional arrays (MDarray type and operators). |

Introduction to the SQL language

- SQL = Structured *Query* Language
 - A standard of the American National Standards Institute (ANSI) in 1986, and of the International Organization for Standardization (ISO) in 1987
 - One of the first commercial languages to utilize the **relational** data model
- A **declarative** *non-procedural* language (4GL) supported by existing RDBMSs
 - Based on tuple relational calculus, associated with relational algebra for *internal* representation, processing, and optimization
 - Specify “**WHAT**” in data requests on a database
 - **NOT** specify “*HOW*” to process the data requests

SQL vs. The Relational Model

| SQL | The Relational Model |
|---|--|
| Table (set if a primary key is specified) | Relation (set) |
| Column | Attribute (+ <i>attribute values</i>) |
| Row | Tuple |
| Domain | Domain |
| Data type | (<i>Data type</i>) |
| Primary key | Primary key |
| Uniqueness, NOT NULL | Secondary key |
| Foreign key | Foreign key |
| Defined and enforced with CHECK, ASSERTION, TRIGGER | Semantic constraints |

SQL vs. The Relational Algebra

| SQL | The Relational Algebra |
|--------------------------------------|---|
| SELECT statement - specifies WHAT | Relational algebra expression - specifies WHAT + HOW |
| FROM clause | (none), x, *, other joins |
| WHERE clause | SELECTION (σ) |
| SELECT clause (+ DISTINCT) | PROJECTION (π) |
| Aggregate functions | Aggregate functions |
| GROUP BY clause | <grouping attributes> \bowtie |
| HAVING clause | SELECTION on \bowtie |
| ORDER BY clause | (not available) |
| UNION, INTERSECT, MINUS | \cup , \cap , - |
| INSERT, DELETE, UPDATE statements | (not available) |

SQL Data Types

- Predefined data types
 - Character Types: Character fixed (CHAR) (e.g. CHAR(10)), Character Varying (VARCHAR) (e.g. VARCHAR(20)), Character Large Object (CLOB)
 - Binary Types: Binary (BINARY), Binary Varying (VARBINARY), Binary Large Object (BLOB)
 - Numeric Types
 - Exact Numeric Types (NUMERIC, DECIMAL, SMALLINT, INTEGER, BIGINT)
 - Approximate Numeric Types (FLOAT, REAL, DOUBLE PRECISION)
 - Datetime Types (DATE, TIME, TIMESTAMP)
 - Interval Type (INTERVAL)
 - Boolean (three-valued logic (3VL)): true, false, *unknown* (NULL)
 - XML, JSON
- Constructed types
 - ARRAY, MULTISET, REF(erence), or ROW
- User-defined types
 - comparable to classes in object-oriented language with their own constructors, observers, mutators, methods, inheritance, overloading, overwriting, interfaces, ...

SQL Data Types – 3-Valued Logic

| AND | True (T) | False (F) | Unknown (U) |
|---------|-------------|--------------|----------------|
| True | T | F | U |
| False | F | F | F |
| Unknown | U | F | U |

| NOT | X |
|---------|---|
| True | F |
| False | T |
| Unknown | U |

| OR | True | False | Unknown |
|---------|------|-------|---------|
| True | T | T | T |
| False | T | F | U |
| Unknown | T | U | U |

Unknown \approx NULL

X: True,
 False,
 Unknown

Introduction to the SQL language

- Data Definition Language (DDL)
 - CREATE, ALTER, DROP
- Data Manipulation Language (DML)
 - For *queries*: SELECT
 - For *data modifications*: INSERT, DELETE, UPDATE
- Data Control Language (DCL)
 - For *access control*: GRANT, REVOKE
 - For *transaction control*: COMMIT, ROLLBACK,
SET AUTOCOMMIT OFF
- SQL: *case-insensitive, extended in DBMSs*

SQL

Simplified SQL Statements (*Checked with each DBMS*)

```
CREATE TABLE <table name> ( <column name><column type> [ <attribute constraint> ]
    { , <column name><column type> [ <attribute constraint> ] }
    [ <table constraint> { , <table constraint> } ] )

DROP TABLE <table name>

ALTER TABLE <table name> ADD <column name><column type>

SELECT [ DISTINCT ] <attribute list>
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }
[ WHERE <condition> ]
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]

<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) )
    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) ) )

<grouping attributes> ::= <column name> { , <column name> }

<order> ::= ( ASC | DESC )

INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }
| <select statement> )

DELETE FROM <table name>
[ WHERE <selection condition> ]

UPDATE <table name>
SET <column name> = <value expression> { , <column name> = <value expression> }
[ WHERE <selection condition> ]

CREATE [ UNIQUE ] INDEX <index name>
ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )
[ CLUSTER ]

DROP INDEX <index name>

CREATE VIEW <view name> [ ( <column name> { , <column name> } ) ]
AS <select statement>

DROP VIEW <view name>
```

Data Definition Language

- Purpose

- Create, modify, and remove the constructs at the structure level of a database

- Named constructs

- Database (schema), tables, types, domains, constraints, indexes, views, assertions, triggers, ...

- DDL statements

- CREATE
 - DROP
 - ALTER

CREATE SCHEMA

- Starting with SQL2, a *schema* is defined to group together tables and other constructs that belong to the *same* database application.
 - tables, types, constraints, views, domains, indexes, authorization grants, etc.
- A *catalog* is a named collection of schemas.
 - Database installations typically have a default environment and schema, so when a user connects and logs in to that database installation, the user can refer directly to tables and other constructs within that schema without having to specify a particular schema name.
 - In a catalog, a special schema is INFORMATION_SCHEMA, which provides information on all the schemas in the catalog and all the element descriptors in these schemas.
 - Schemas within the same catalog can also share certain elements, such as type and domain definitions.

CREATE SCHEMA

- An SQL schema is identified by a *schema name* and includes an *authorization identifier* to indicate the user or account who owns the schema, and descriptors for each element.
- The *privilege* to create schemas, tables, and other constructs must be explicitly granted to the relevant *user* accounts by the system administrator or DBA.

`CREATE SCHEMA SchemaName
[AUTHORIZATION AuthorizationIdentifier];`

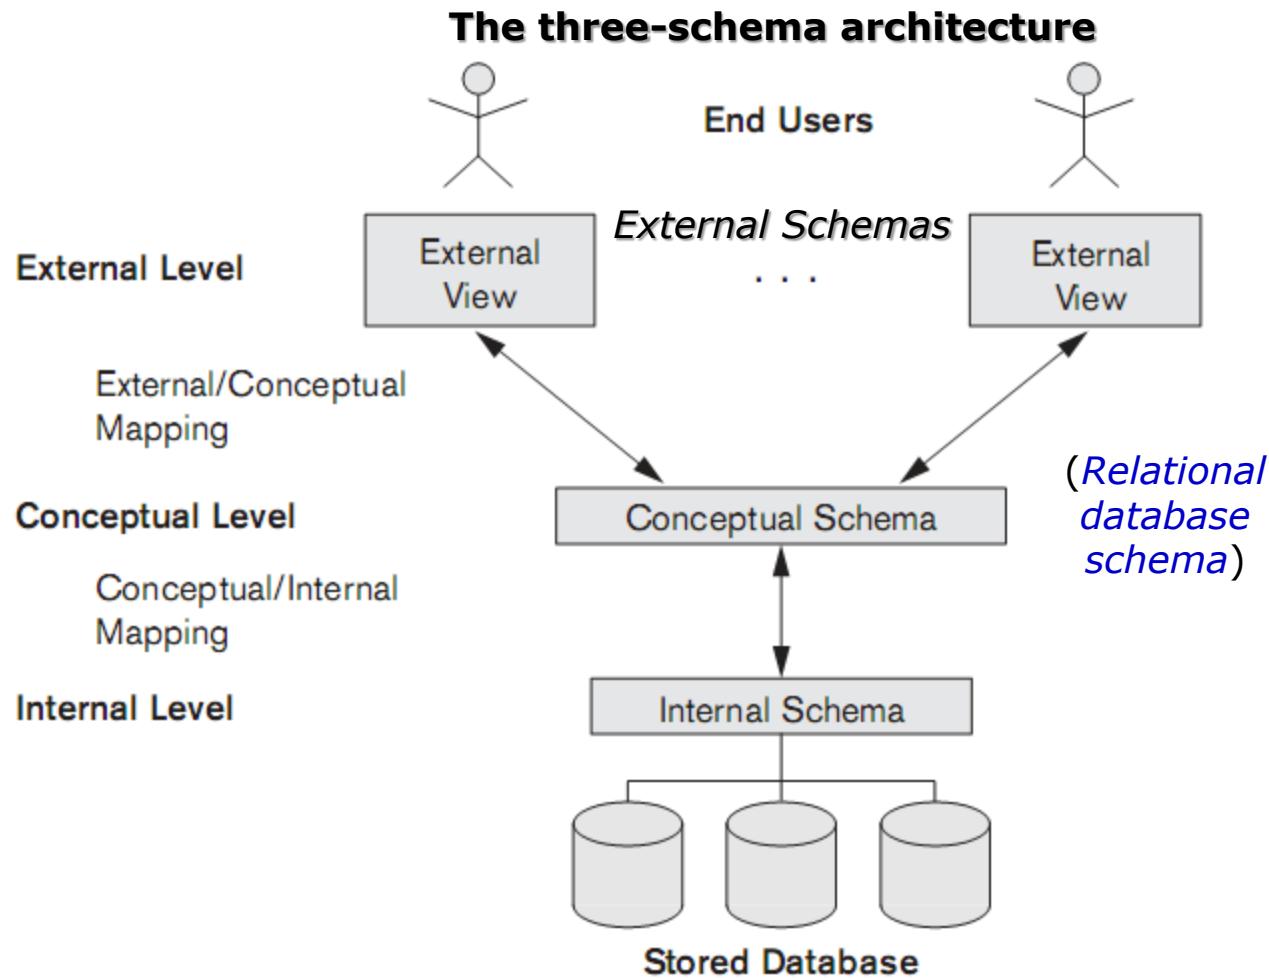
`CREATE SCHEMA Company AUTHORIZATION Smith;
CREATE SCHEMA Company;`

NOTES on SCHEMAS

the part of the database that a particular user group is interested in and hides the rest of the database from that user group

the structure of the whole database for a community of users

the physical storage structure of the database



- *Conceptual schema in the three-schema architecture ≠ the conceptual EER schema in design!!!*
- Conceptual schema in the three-schema architecture is in fact *the relational database schema*.
- External Schemas are supported with CREATE VIEW.

NOTES on SCHEMAS

- MySQL
 - CREATE SCHEMA = CREATE DATABASE
 - Concepts are equivalent to the standard.
- MS SQL Server
 - CREATE DATABASE with more physical properties
 - Data/ log files
 - **dbo** (the default schema of the database owner) is a default schema for each database to group objects of similar scope or ownership.
- Oracle
 - CREATE DATABASE with more physical properties
 - Data/ log files
 - CREATE USER with a schema of the same name created by default
 - A user is an account through which we can log in to the database.
 - CREATE SCHEMA does not actually create a schema. Oracle Database automatically creates a schema when a user is created. This statement populates the schema (*the one with the username*) with tables and views and grant privileges on those objects without having to issue multiple SQL statements in multiple transactions.

CREATE TABLE

- A base table is created and actually stored as a file by a DBMS with **CREATE TABLE**.
 - The attributes in a base table are considered to be ordered in the sequence in which they are specified in the CREATE TABLE statement.
 - Rows (tuples) are not considered to be ordered within a table (relation).
- A virtual table is created but may or may not correspond to an actual physical file with **CREATE VIEW**.

CREATE TABLE

```
CREATE TABLE [SchemaName.]TableName  
({columnName dataType [NOT NULL] [UNIQUE] [PRIMARY KEY]  
    [DEFAULT <value>]  
    [CHECK searchCondition] [,...]}  
[PRIMARY KEY (listOfColumns)] [,...]  
{[UNIQUE (listOfColumns)] [,...]}  
{[FOREIGN KEY (listOfFKColumns)  
    REFERENCES [SchemaName.] ParentTableName [(listOfPKColumns)]  
    [ON UPDATE CASCADE | SET NULL]  
    [ON DELETE CASCADE | SET NULL] [,...]}  
{[CHECK (searchCondition)] [,...]}  
{[CONSTRAINT constraintName PRIMARY KEY| UNIQUE| FOREIGN KEY|  
    CHECK constraint...] [,...]})
```

The COMPANY database

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|-----|-------|---------|-----|--------|-----------|-----|
|-------|-------|-------|-----|-------|---------|-----|--------|-----------|-----|



DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|-------|---------|---------|----------------|
|-------|---------|---------|----------------|



DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
|---------|-----------|

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-------|---------|-----------|------|
|-------|---------|-----------|------|



WORKS_ON

| Essn | Pno | Hours |
|------|-----|-------|
|------|-----|-------|

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|------|----------------|-----|-------|--------------|
|------|----------------|-----|-------|--------------|

```
CREATE SCHEMA COMPANY;
```

```
CREATE TABLE EMPLOYEE (
    FNAME            VARCHAR(15)      NOT NULL,
    MINIT           CHAR,
    LNAME            VARCHAR(15)      NOT NULL,
    SSN              CHAR(9),
    BDATE            DATE,
    ADDRESS          VARCHAR(30),
    SEX              CHAR,
    SALARY           DECIMAL(10,2),
    SUPER_SSN        CHAR(9),
    DNO              INT             NOT NULL,
    PRIMARY KEY (SSN),
    FOREIGN KEY (SUPER_SSN) REFERENCES EMPLOYEE (SSN)
);
```

```
CREATE TABLE DEPENDENT (
    ESSN             CHAR(9),
    DEPENDENT_NAME   VARCHAR(15),
    SEX              CHAR,
    BDATE            DATE,
    RELATIONSHIP     VARCHAR(8),
    PRIMARY KEY (ESSN, DEPENDENT_NAME),
    FOREIGN KEY (ESSN) REFERENCES EMPLOYEE (SSN)
);
```

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(15)      NOT NULL,
    DNUMBER        INT,
    MGR_SSN        CHAR(9)         NOT NULL,
    MGR_STAR_TDATE DATE,
    PRIMARY KEY (DNUMBER),
    UNIQUE (DNAME),
    FOREIGN KEY (MGR_SSN) REFERENCES EMPLOYEE (SSN)
);
```

```
CREATE TABLE DEPT_LOCATIONS (
    DNUMBER        INT,
    DLOCATION      VARCHAR(15),
    PRIMARY KEY (DNUMBER, DLOCATION),
    FOREIGN KEY (DNUMBER) REFERENCES DEPARTMENT(DNUMBER)
);
```

```
ALTER TABLE EMPLOYEE
ADD FOREIGN KEY (DNO) REFERENCES DEPARTMENT (DNUMBER);
```

```
CREATE TABLE PROJECT (
    PNAME          VARCHAR(15)      NOT NULL,
    PNUMBER        INT,
    PLOCATION      VARCHAR(15),
    DNUM           INT             NOT NULL,
    PRIMARY KEY (PNUMBER),
    UNIQUE (PNAME),
    FOREIGN KEY (DNUM) REFERENCES DEPARTMENT (DNUMBER)
);

CREATE TABLE WORKS_ON (
    ESSN          CHAR(9),
    PNO           INT,
    HOURS         DECIMAL(3,1),
    PRIMARY KEY (ESSN, PNO),
    FOREIGN KEY (ESSN) REFERENCES EMPLOYEE (SSN),
    FOREIGN KEY (PNO) REFERENCES PROJECT (PNUMBER)
);
```

CREATE TABLE - CONSTRAINTS

- Basic constraints are specified as part of table creation.
 - Primary key
 - Foreign key (Referential integrity constraint)
 - UNIQUE
 - NOT NULL
 - Attribute defaults
 - Restrictions on attribute domains and other constraints on individual tuples within a table using the CHECK clause

CREATE TABLE - CONSTRAINTS

- Primary key: specified *only one* for each table

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(15),
    DNUMBER        INT          PRIMARY KEY,
    MGRSSN         CHAR(9),
    MGRSTARTDATE   DATE
);
```

```
CREATE TABLE WORKS_ON (
    ESSN          CHAR(9),
    PNO           INT,
    HOURS         DECIMAL(3,1),
    PRIMARY KEY (ESSN, PNO)
);
```

CREATE TABLE - CONSTRAINTS

- Foreign key: specified for *referential integrity*

```
CREATE TABLE DEPARTMENT (
```

| | |
|--------------|--------------------------|
| DNAME | VARCHAR(15), |
| DNUMBER | INT PRIMARY KEY , |
| MGRSSN | CHAR(9), |
| MGRSTARTDATE | DATE |

```
);
```

```
CREATE TABLE DEPT_LOCATIONS (
```

| | |
|---|--------------|
| DNUMBER | INT, |
| DLOCATION | VARCHAR(15), |
| PRIMARY KEY (DNUMBER, DLOCATION), | |
| FOREIGN KEY (DNUMBER) REFERENCES DEPARTMENT(DNUMBER) | |

```
);
```

CREATE TABLE - CONSTRAINTS

- Foreign key: specified for *referential integrity*
 - The default action for an integrity violation is to *reject* the *update* operation that will cause a violation: RESTRICT.
 - Other actions: SET NULL, CASCADE, and SET DEFAULT
 - CASCADE ON DELETE is to delete all the referencing tuples, whereas CASCADE ON UPDATE is to change the value of the referencing foreign key attribute(s) to the updated (new) primary key value for all the referencing tuples.

```
CREATE TABLE DEPT_LOCATIONS (
    DNUMBER      INT,
    DLOCATION    VARCHAR(15),
    PRIMARY KEY (DNUMBER, DLOCATION),
    FOREIGN KEY (DNUMBER) REFERENCES DEPARTMENT(DNUMBER)
                           ON DELETE CASCADE ON UPDATE CASCADE
);
```

CREATE TABLE - CONSTRAINTS

- ❑ UNIQUE: no duplicate values for the attribute
- ❑ NOT NULL: no NULL values allowed. NULL is default with no specification.

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(15)  UNIQUE,
    DNUMBER        INT          PRIMARY KEY,
    MGRSSN         CHAR(9)      NOT NULL,
    MGRSTARTDATE   DATE
);
```

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(15), .....
    UNIQUE (DNAME)
);
```

CREATE TABLE - CONSTRAINTS

- Attribute defaults: a default value for an attribute by appending DEFAULT <value>
 - The default value is included in any new tuple if an explicit value is not provided for that attribute.
 - If no default clause is specified, the default value is NULL for attributes that do not have the NOT NULL constraint.

```
CREATE TABLE WORKS_ON (
    ESSN           CHAR(9),
    PNO            INT,
    HOURS          DECIMAL(3,1) DEFAULT 10.0,
    PRIMARY KEY (ESSN, PNO)
);
```

CREATE TABLE - CONSTRAINTS

- Restrictions on attribute domains and other constraints on individual tuples within a table using the CHECK clause: *row-level constraints*

Constraint: Department numbers are restricted to integer numbers between 1 and 20. Starting date is from 2000-01-01.

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(15)  UNIQUE,
    DNUMBER        INT          PRIMARY KEY
                            CHECK (DNUMBER>0 AND DNUMBER<21),
    MGRSSN         CHAR(9)      NOT NULL,
    MGRSTARTDATE   DATE
                            CHECK (MGRSTARTDATE> '2000-01-01')
);
```

CREATE DOMAIN

- ❑ A domain can be declared, and the domain name can be used with the attribute specification. CREATE DOMAIN is *DBMS-specifically* implemented.
 - change the data type for a domain that is used by numerous attributes in a schema, and improve schema *readability*

CREATE DOMAIN *DomainName* AS *DataType*
[**CHECK** (<*condition*>)];

```
CREATE DOMAIN SSN_TYPE AS CHAR(9);
```

```
CREATE DOMAIN D_NUM AS INTEGER  
    CHECK (D_NUM > 0 AND D_NUM < 21);
```

CREATE DOMAIN

```
CREATE DOMAIN SSN_TYPE AS CHAR(9);
```

```
CREATE DOMAIN D_NUM AS INTEGER  
    CHECK (D_NUM > 0 AND D_NUM < 21);
```

```
CREATE TABLE DEPARTMENT (  
    DNAME          VARCHAR(15)  UNIQUE,  
    DNUMBER        D_NUM        PRIMARY KEY,  
    MGRSSN         SSN_TYPE     NOT NULL,  
    MGRSTARTDATE   DATE  
);
```

What is difference between CHECK at the attribute level and CHECK at the domain level?

CREATE TABLE - CONSTRAINTS

□ Giving *names* to constraints: CONSTRAINT

- This is optional.
- The name is unique within a particular database schema.
- It is used to identify a particular constraint in case it must be dropped later and replaced with another one.
- It is also possible to temporarily defer a constraint until the end of a transaction.

CREATE TABLE EMPLOYEE

```
( ...,
    DNO          INT  NOT NULL  DEFAULT 1,
CONSTRAINT EMPPK
    PRIMARY KEY (SSN) ,
CONSTRAINT EMPSUPERFK
    FOREIGN KEY (SUPERSSN) REFERENCES EMPLOYEE(SSN)
        ON DELETE SET NULL  ON UPDATE CASCADE ,
CONSTRAINT EMPDEPTFK
    FOREIGN KEY (DNO) REFERENCES DEPARTMENT(DNUMBER)
        ON DELETE SET DEFAULT  ON UPDATE CASCADE );
```

CREATE TABLE DEPARTMENT

```
( ...,
    MGRSSN  CHAR(9) NOT NULL DEFAULT '888665555' ,
    ...,
CONSTRAINT DEPTPK
    PRIMARY KEY (DNUMBER) ,
CONSTRAINT DEPTSK
    UNIQUE (DNAME),
CONSTRAINT DEPTMGRFK
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE(SSN)
        ON DELETE SET DEFAULT  ON UPDATE CASCADE );
```

CREATE TABLE DEPT_LOCATIONS

```
( ...,
    PRIMARY KEY (DNUMBER, DLOCATION),
    FOREIGN KEY (DNUMBER) REFERENCES DEPARTMENT(DNUMBER)
        ON DELETE CASCADE  ON UPDATE CASCADE );
```

DROP Statement

- ❑ Used to drop named schema elements: schema (*database*), tables, indexes, functions, triggers, ...
- ❑ Two drop behavior options:
 - CASCADE: all the elements referencing or contained inside the one being dropped are also dropped.
 - RESTRICT: only the one being dropped if no reference.
- ❑ **DROP SCHEMA**

DROP SCHEMA *Company* CASCADE;

DROP SCHEMA *Company* RESTRICT;

DROP Statement

□ DROP TABLE:

DROP TABLE *Dependent CASCADE*;

- CASCADE: all such constraints, views, and other elements that reference the table are dropped automatically from the schema along with the table itself.

DROP TABLE *Dependent RESTRICT*;

- RESTRICT: dropped if it is not referenced in any constraints, views, and other elements.

DROP TABLE DEPARTMENT RESTRICT;

**Error Code: 1217. Cannot delete or update a parent row:
a foreign key constraint fails (MySQL)**

ALTER Statement

- Modifications at the structure level:

- ALTER DATABASE
- ALTER TABLE
- ALTER VIEW
- ALTER FUNCTION
- ALTER PROCEDURE
- ...

- ALTER TABLE:

- COLUMN: ADD, MODIFY, DROP, ...
- CONSTRAINT: ADD, DROP, ...
- TABLE PROPERTIES
- ...

ALTER Statement

- Add an attribute for keeping track of jobs of employees to EMPLOYEE table in COMPANY schema

ALTER TABLE Company.Employee ADD (Job VARCHAR(12));

- Job value for each existing row: default value

ALTER TABLE Company.Employee ADD (Job VARCHAR(12) DEFAULT 'EMPLOYEE');

- What value does each existing row take for attribute Job if?

ALTER TABLE Company.Employee ADD (Job VARCHAR(12) NOT NULL);

- What happens to *two or more existing rows* if attribute Job is *not null and unique*?

ALTER TABLE Company.Employee ADD (Job VARCHAR(12) **NOT NULL UNIQUE**);

ALTER Statement

- Drop a column: similarly to drop a table, CASCADE or RESTRICT option must be specified
 - CASCADE option: all constraints and views referencing the column are dropped along with the column
 - RESTRICT option: successful only if no constraints and views are referencing the column

```
ALTER TABLE Company.Employee DROP [COLUMN] Address  
CASCADE;
```

- What if a *primary key* is dropped?

```
ALTER TABLE Company.Employee DROP Ssn CASCADE;
```

ALTER Statement

- Add a constraint to a table. Added constraints must be satisfied by the current content.

- Primary key

```
ALTER TABLE Works_on ADD PRIMARY KEY (Ssn, Pnumber);
```

- Foreign key

```
ALTER TABLE Works_on ADD CONSTRAINT ssn_fk  
FOREIGN KEY (Ssn) REFERENCES Employee (Ssn)  
ON DELETE CASCADE ON UPDATE CASCADE;
```

- Default values

```
ALTER TABLE Employee ALTER COLUMN Dno SET DEFAULT 1;
```

- Unique

```
ALTER TABLE Department ADD UNIQUE (Dname);
```

Data Definition Language

□ Drop a constraint

```
ALTER TABLE Works_on  
DROP PRIMARY KEY;
```

```
ALTER TABLE Works_on  
DROP FOREIGN KEY ssn_fk CASCADE;
```

```
ALTER TABLE Employee  
ALTER COLUMN Dno DROP DEFAULT;
```

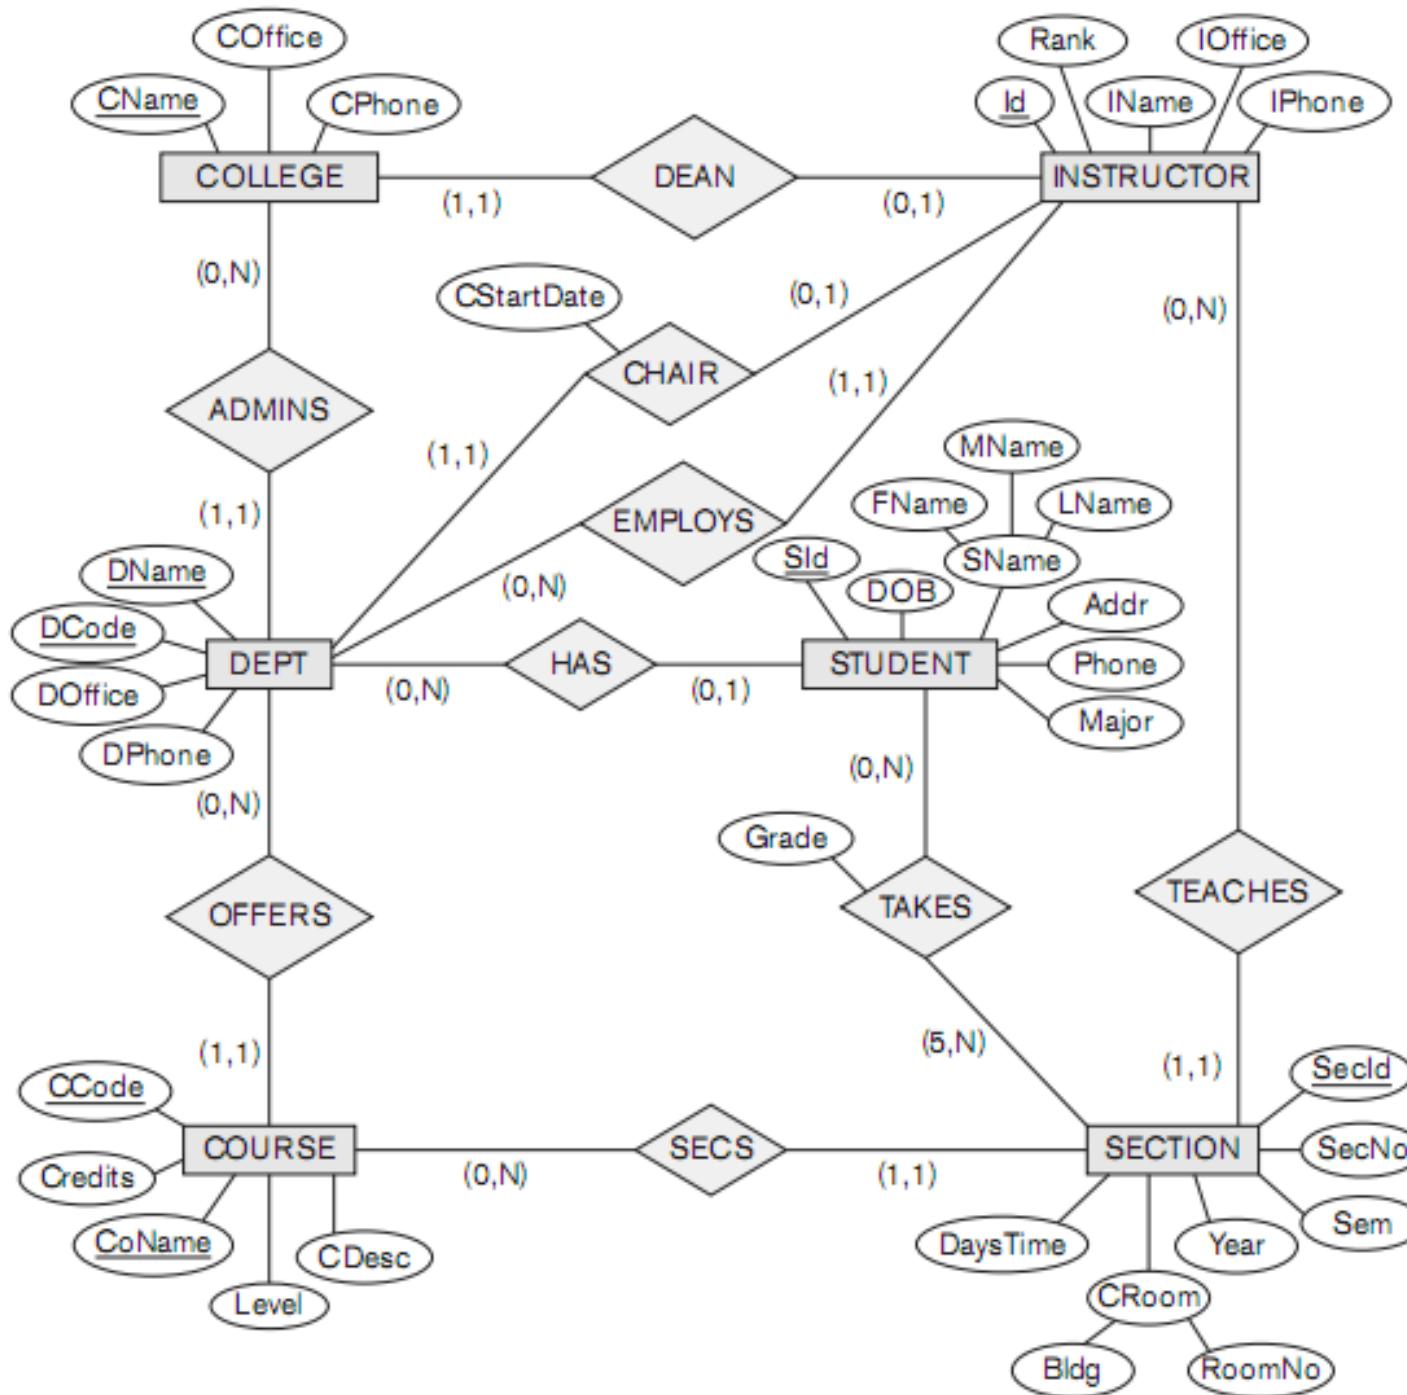
//Oracle

```
ALTER TABLE Department  
DROP UNIQUE (Dname);
```

//MySQL – *DROP UNIQUE*

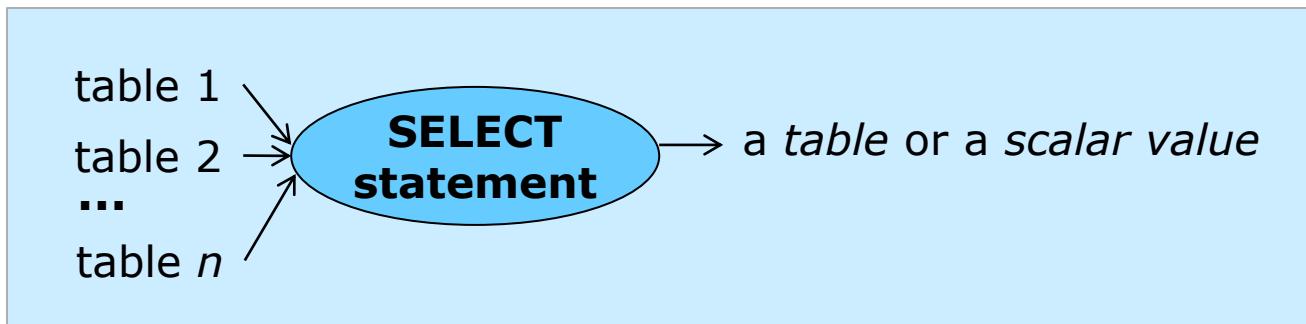
```
ALTER TABLE Department  
DROP INDEX Dname;
```

Write DDL statements for the UNIVERSITY database on a *specific* DBMS. Pay attention to the references among the tables and their *other* constraints.



SELECT Statement

- SQL has one basic statement for retrieving information from a database: SELECT statement.
- SELECT statement is *not the same as* SELECTION operation (σ) of the relational algebra.
- Important distinction between the SQL data model and the formal relational model is that: SQL allows a table to have two or more tuples that are identical in all their attribute values.
 - A SQL table is a *multi-set* (sometimes called a bag) of tuples, which *is not* a set of tuples.
 - A SQL table is a set of tuples, like a relation in the relational model, once a primary key is defined on it.



SELECT Statement

SELECT [DISTINCT | ALL]

{* | [columnExpression [AS newName]] | built-inFunction
[,...] }

[**FROM** TableName|ViewName [alias] [, ...]]

[**WHERE** rowCondition]

[**GROUP BY** columnList]

[**HAVING** groupCondition]

[**ORDER BY** columnList | columnPositionList [ASC|DESC]]

SELECT Statement

Typical processing of a query block

| Execution order | Clause | Meaning |
|-----------------|---------------|--|
| 1 | FROM | Specifies <i>and joins</i> table(s)/ view(s) to be used |
| 2 | WHERE | Filters rows by row-level conditions |
| 3 | GROUP BY | Forms groups of rows with the same grouping column value |
| 4 | HAVING | Filters groups by group-level conditions |
| 5 | SELECT | Specifies the output to be returned |
| 6 | ORDER BY | Specifies the order of the output |

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

The COMPANY database

DEPARTMENT

| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|----------------|----------------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

| <u>Pname</u> | <u>Rnumber</u> | <u>Plocation</u> | <u>Dnum</u> |
|-----------------|----------------|------------------|-------------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |

Source: [1]

SELECT Statement

For each project in Stafford, *on which more than two employees work*, retrieve the project name and the number of employees who work on that project in descending order of the project name.

```
SELECT          PNAME, COUNT (*)
FROM           PROJECT JOIN WORKS_ON ON PNUMBER=PNO
WHERE          PLOCATION = 'Stafford'
GROUP BY        PNAME
HAVING         COUNT (*) > 2
ORDER BY        PNAME DESC;
```

SELECT Statement

FROM PROJECT JOIN WORKS_ON ON PNUMBER=PNO

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

⋈ PNUMBER=PNO

WORKS_ON

| Essn | Pho | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

SELECT Statement

FROM PROJECT JOIN WORKS_ON ON PNUMBER=PNO

| Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|-----------------|---------|-----------|------|-----------|-----|-------|
| ProductX | 1 | Bellaire | 5 | 123456789 | 1 | 32.5 |
| ProductX | 1 | Bellaire | 5 | 453453453 | 1 | 20.0 |
| ProductY | 2 | Sugarland | 5 | 123456789 | 2 | 7.5 |
| ProductY | 2 | Sugarland | 5 | 453453453 | 2 | 20.0 |
| ProductY | 2 | Sugarland | 5 | 333445555 | 2 | 10.0 |
| ProductZ | 3 | Houston | 5 | 666884444 | 3 | 40.0 |
| ProductZ | 3 | Houston | 5 | 333445555 | 3 | 10.0 |
| Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Reorganization | 20 | Houston | 1 | 333445555 | 20 | 10.0 |
| Reorganization | 20 | Houston | 1 | 987654321 | 20 | 15.0 |
| Reorganization | 20 | Houston | 1 | 888665555 | 20 | NULL |
| Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

SELECT Statement

WHERE PLOCATION = 'Stafford'

| Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|-----------------|---------|-----------|------|-----------|-----|-------|
| ProductX | 1 | Bellaire | 5 | 123456789 | 1 | 32.5 |
| ProductX | 1 | Bellaire | 5 | 453453453 | 1 | 20.0 |
| ProductY | 2 | Sugarland | 5 | 123456789 | 2 | 7.5 |
| ProductY | 2 | Sugarland | 5 | 453453453 | 2 | 20.0 |
| ProductY | 2 | Sugarland | 5 | 333445555 | 2 | 10.0 |
| ProductZ | 3 | Houston | 5 | 666884444 | 3 | 40.0 |
| ProductZ | 3 | Houston | 5 | 333445555 | 3 | 10.0 |
| Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Reorganization | 20 | Houston | 1 | 333445555 | 20 | 10.0 |
| Reorganization | 20 | Houston | 1 | 987654321 | 20 | 15.0 |
| Reorganization | 20 | Houston | 1 | 888665555 | 20 | NULL |
| Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

SELECT Statement

WHERE PLOCATION = 'Stafford'

| Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|-----------------|---------|-----------|------|-----------|-----|-------|
| Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

SELECT Statement

GROUP BY PNAME

| | Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|---------|-----------------|---------|-----------|------|-----------|-----|-------|
| Group 1 | Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Group 2 | Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| | Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| | Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

SELECT Statement

HAVING COUNT (*) > 2

| | Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|---------|-----------------|----------------|------------------|-------------|-------------|------------|--------------|
| Group 1 | Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Group 2 | Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| | Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| | Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

Group 1: COUNT(*) = 3 > 2 => TRUE => group 1 selected

Group 2: COUNT(*) = 3 > 2 => TRUE => group 2 selected

SELECT Statement

SELECT PNAME, COUNT (*)

| | Pname | Pnumber | Plocation | Dnum | Essn | Pno | Hours |
|---------|-----------------|----------------|------------------|-------------|-------------|------------|--------------|
| Group 1 | Computerization | 10 | Stafford | 4 | 333445555 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 999887777 | 10 | 10.0 |
| | Computerization | 10 | Stafford | 4 | 987987987 | 10 | 35.0 |
| Group 2 | Newbenefits | 30 | Stafford | 4 | 999887777 | 30 | 30.0 |
| | Newbenefits | 30 | Stafford | 4 | 987987987 | 30 | 5.0 |
| | Newbenefits | 30 | Stafford | 4 | 987654321 | 30 | 20.0 |

Group 1: COUNT(*) = 3 > 2 => TRUE => group 1 selected

Group 2: COUNT(*) = 3 > 2 => TRUE => group 2 selected

| Pname | COUNT(*) |
|-----------------|-----------------|
| Computerization | 3 |
| Newbenefits | 3 |

SELECT Statement

ORDER BY PNAME DESC;

| Pname | COUNT(*) |
|-----------------|----------|
| Computerization | 3 |
| Newbenefits | 3 |

Descending
order



| Pname | COUNT(*) |
|-----------------|----------|
| Newbenefits | 3 |
| Computerization | 3 |

For each project in Stafford, *on which more than two employees work*, retrieve the project name and the number of employees who work on that project in descending order of the project name.

```
SELECT PNAME, COUNT (*)
FROM PROJECT JOIN WORKS_ON ON PNUMBER=PNO
WHERE PLOCATION = 'Stafford'
GROUP BY PNAME
HAVING COUNT (*) > 2
ORDER BY PNAME DESC;
```

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

WORKS_ON

| Esn | Pho | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

Input Tables

Query Result

| Pname | COUNT(*) |
|-----------------|----------|
| Newbenefits | 3 |
| Computerization | 3 |

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

Retrieve all the information of each employee who is with department 4 and salary > 25000 or with department 5 and salary > 30000.

$\sigma(Dno=4 \text{ AND } Salary>25000) \text{ OR } (Dno=5 \text{ AND } Salary>30000)$ (EMPLOYEE)

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |

SELECT *

FROM EMPLOYEE

WHERE (Dno = 4 **AND** Salary > 25000) **OR** (Dno = 5 **AND** SALARY > 30000);

* (asterisk): all columns in the previous process are included in the output.

The primary key in the output => the output is a set (relation).

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | | | | |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | | | | |

Retrieve Ssn, date of birth, and department number of each employee.

$\pi_{\text{Ssn}, \text{Bdate}, \text{Dno}}(\text{EMPLOYEE})$

SELECT Ssn, Bdate, Dno

FROM EMPLOYEE;

| Ssn | Bdate | Dno |
|-----------|------------|-----|
| 123456789 | 1965-01-09 | 5 |
| 333445555 | 1955-12-08 | 5 |
| 999887777 | 1968-01-19 | 4 |
| 987654321 | 1941-06-20 | 4 |
| 666884444 | 1962-09-15 | 5 |
| 453453453 | 1972-07-31 | 5 |
| 987987987 | 1969-03-29 | 4 |
| 888665555 | 1937-11-10 | 1 |

List all the department numbers of the employees.

$\pi_{\text{Dno}}(\text{EMPLOYEE})$

SELECT DISTINCT Dno

FROM EMPLOYEE;

| Dno |
|-----|
| 5 |
| 4 |
| 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

Return all the combinations of departments and their locations.

DEPARTMENT x DEPT_LOCATIONS

| Dname | Dnumber | Mgr_ssn | Mgr_start_date | Dnumber | Dlocation |
|----------------|---------|-----------|----------------|---------|-----------|
| Research | 5 | 333445555 | 1988-05-22 | 1 | Houston |
| Research | 5 | 333445555 | 1988-05-22 | 4 | Stafford |
| Research | 5 | 333445555 | 1988-05-22 | 5 | Bellaire |
| Research | 5 | 333445555 | 1988-05-22 | 5 | Sugarland |
| Administration | 4 | 987654321 | 1995-01-01 | 1 | Houston |
| Administration | 4 | 987654321 | 1995-01-01 | 4 | Stafford |
| Administration | 4 | 987654321 | 1995-01-01 | 5 | Bellaire |
| Administration | 4 | 987654321 | 1995-01-01 | 5 | Sugarland |
| Headquarters | 1 | 888665555 | 1981-06-19 | 1 | Houston |
| Headquarters | 1 | 888665555 | 1981-06-19 | 4 | Stafford |
| Headquarters | 1 | 888665555 | 1981-06-19 | 5 | Bellaire |
| Headquarters | 1 | 888665555 | 1981-06-19 | 5 | Sugarland |
| Headquarters | 1 | 888665555 | 1981-06-19 | 5 | Houston |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

Return all the combinations of Research department with Houston location.

DEPARTMENT  **Dname = 'Research' AND Dlocation = 'Houston' DEPT_LOCATIONS**

| Dname | Dnumber | Mgr_ssn | Mgr_start_date | Dnumber | Dlocation |
|----------|---------|-----------|----------------|---------|-----------|
| Research | 5 | 333445555 | 1988-05-22 | 1 | Houston |
| Research | 5 | 333445555 | 1988-05-22 | 5 | Houston |

SELECT *

FROM DEPARTMENT, DEPT_LOCATIONS;

WHERE Dname = 'Research' AND Dlocation = 'Houston';

SELECT *

FROM DEPARTMENT **CROSS JOIN DEPT_LOCATIONS;**

WHERE Dname = 'Research' AND Dlocation = 'Houston';

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

Retrieve all the information of each department and its locations.

DEPARTMENT  **Dnumber = Dnumber DEPT_LOCATIONS**

| Dname | Dnumber | Mgr_ssn | Mgr_start_date | Dnumber | Dlocation |
|----------------|---------|-----------|----------------|---------|-----------|
| Research | 5 | 333445555 | 1988-05-22 | 5 | Bellaire |
| Research | 5 | 333445555 | 1988-05-22 | 5 | Sugarland |
| Research | 5 | 333445555 | 1988-05-22 | 5 | Houston |
| Administration | 4 | 987654321 | 1995-01-01 | 4 | Stafford |
| Headquarters | 1 | 888665555 | 1981-06-19 | 1 | Houston |

```
SELECT *
FROM DEPARTMENT d, DEPT_LOCATIONS dl;
WHERE d.Dnumber = dl.Dnumber;
```

```
SELECT *
FROM DEPARTMENT d CROSS JOIN DEPT_LOCATIONS dl;
WHERE d.Dnumber = dl.Dnumber;
```

```
SELECT *
FROM DEPARTMENT d JOIN DEPT_LOCATIONS dl ON d.Dnumber = dl.Dnumber;
```

Ambiguous names => qualified names by [schemaName.]tableName.attributeName
ALIAS => alternative table name => tuple variable for the table

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

Retrieve all the information of each department and its locations.

DEPARTMENT * DEPT_LOCATIONS

| Dname | Dnumber | Mgr_ssn | Mgr_start_date | Dlocation |
|----------------|---------|-----------|----------------|-----------|
| Research | 5 | 333445555 | 1988-05-22 | Bellaire |
| Research | 5 | 333445555 | 1988-05-22 | Sugarland |
| Research | 5 | 333445555 | 1988-05-22 | Houston |
| Administration | 4 | 987654321 | 1995-01-01 | Stafford |
| Headquarters | 1 | 888665555 | 1981-06-19 | Houston |

SELECT *

FROM DEPARTMENT NATURAL JOIN DEPT_LOCATIONS;

Return the information of all the employees and their departments they manage.

EMPLOYEE **Ssn = Mgr_ssn** DEPARTMENT

| Fname | Minit | Lname | Ssn | B_date | Address | Sex | Salary | Super_ssn | Dno | Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------|-------|---------|---------------|--------------------|-----------------------------------|-----|--------|---------------|-----|--------------------|---------|---------------|----------------|
| Franklin | T | Wong | 33344 5555 | 1955 -12- 08 | 638 Voss, Houst on, TX | M | 40000 | 888665 555 | 5 | Resear ch | 5 | 33344 5555 | 1988- 05-22 |
| Jennifer | S | Wallace | 98765 4321 | 1941 -06- 20 | 291 Berry, Bellair e, TX | F | 43000 | 888665 555 | 4 | Admini stration | 4 | 98765 4321 | 1995- 01-01 |
| James | E | Borg | 88866 5555 | 1937 -11- 10 | 450 Stone, Houst on, TX | M | 55000 | NULL | 1 | Headqu arters | 1 | 88866 5555 | 1981- 06-09 |

```
SELECT *
FROM EMPLOYEE, DEPARTMENT;
WHERE Ssn = Mgr_ssn;
```

```
SELECT *
FROM EMPLOYEE CROSS JOIN DEPARTMENT;
WHERE Ssn = Mgr_ssn;
```

```
SELECT *
FROM EMPLOYEE JOIN DEPARTMENT ON Ssn = Mgr_ssn;
```

Return the information of all the employees and their departments that they manage if any.

EMPLOYEE  **Ssn = Mgr_ssn DEPARTMENT**

| Fname | Minit | Lname | Ssn | B date | Address | Sex | Salary | Super_ssn | Dno | Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------|-------|---------|-------------------|--------------------|--|-----|--------|---------------|-----|--------------------|---------|-------------------|----------------|
| Franklin | T | Wong | 3334 4555 5 | 1955 -12- 08 | 638 Voss, Houst on, TX | M | 40000 | 888665 555 | 5 | Resear ch | 5 | 3334 4555 5 | 1988- 05-22 |
| Jennifer | S | Wallace | 9876 5432 1 | 1941 -06- 20 | 291 Berry, Bellair e, TX | F | 43000 | 888665 555 | 4 | Admini stration | 4 | 9876 5432 1 | 1995- 01-01 |
| James | E | Borg | 8886 6555 5 | 1937 -11- 10 | 450 Stone, Houst on, TX | M | 55000 | NULL | 1 | Headqu arters | 1 | 8886 6555 5 | 1981- 06-09 |
| John | B | Smith | 1234 5678 9 | 1965 -01- 09 | 732 Fondre n, Houst on, TX | M | 30000 | 333445 555 | 5 | NULL | NULL | NULL | NULL |
| Alicia | J | Zelaya | 9997 7888 8 | ... | ... | ... | ... | ... | 4 | NULL | NULL | NULL | NULL |
| Ramesh | K | Narayan | ... | ... | ... | ... | ... | ... | 5 | NULL | NULL | NULL | NULL |
| Joyce | A | English | ... | ... | ... | ... | ... | ... | 5 | NULL | NULL | NULL | NULL |
| Ahmad | V | Jabbar | ... | ... | ... | ... | ... | ... | 4 | NULL | NULL | NULL | NULL |

SELECT *

FROM EMPLOYEE LEFT OUTER JOIN DEPARTMENT ON Ssn = Mgr_ssn;

□ SEMI-JOIN: $T_1 \bowtie_{T_1.X=T_2.Y} T_2$

Return the information of all the departments *with at least one* employee who has salary > 30000.

DEPARTMENT $\bowtie_{Dnumber=Dno}$ ($\sigma_{Salary>30000}$ (**EMPLOYEE**))

SELECT Dname, Dnumber, Mgr_ssn, Mgr_start_date

FROM DEPARTMENT **JOIN** (**SELECT** * **FROM** EMPLOYEE **WHERE** Salary>30000) e

ON Dnumber = Dno;

| DNAME | DNUMBER | MGR_SSN | MGR_START_DATE |
|----------------|---------|-----------|----------------|
| Headquarters | 1 | 888665555 | 1981-06-19 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Research | 5 | 333445555 | 1988-05-22 |

SELECT *

FROM DEPARTMENT

WHERE Dnumber **IN** (**SELECT** Dno **FROM** EMPLOYEE **WHERE** Salary > 30000);

SELECT *

FROM DEPARTMENT d

WHERE EXISTS

(**SELECT** * **FROM** EMPLOYEE **WHERE** Dno = d.Dnumber **AND** Salary > 30000);

SEMI-JOIN
(Oracle, MySQL, ...)
(DBMS-specific implementation)

EXISTS vs. NOT EXISTS

- EXISTS function returns TRUE or FALSE.
 - TRUE: if the nested query result contains *at least one tuple*
 - FALSE: if the nested query result contains *no tuples*

Retrieve Ssn and name of each employee who has *no* dependents.

```
SELECT Ssn, Fname, Lname  
FROM EMPLOYEE  
WHERE NOT EXISTS (SELECT *  
                      FROM DEPENDENT  
                      WHERE Essn = Ssn);
```

```
SELECT Ssn, Fname, Lname  
FROM EMPLOYEE  
WHERE Ssn NOT IN (SELECT DISTINCT ESsn  
                      FROM DEPENDENT);
```

| Ssn | Fname | Lname |
|-----------|--------|---------|
| 453453453 | Joyce | English |
| 666884444 | Ramesh | Narayan |
| 888665555 | James | Borg |
| 987987987 | Ahmad | Jabbar |
| 999887777 | Alicia | Zelaya |

EXISTS vs. NOT EXISTS

- EXISTS function returns TRUE or FALSE.
 - TRUE: if the nested query result contains *at least one* tuple
 - FALSE: if the nested query result contains *no* tuples

List Ssn and name of each manager who has *at least one* dependent.

```
SELECT Ssn, Fname, Lname  
FROM EMPLOYEE  
WHERE EXISTS (SELECT *  
                 FROM DEPENDENT  
                 WHERE Essn = Ssn)  
AND EXISTS (SELECT *  
                 FROM DEPARTMENT  
                 WHERE Mgr_ssn = Ssn);
```

| Ssn | Fname | Lname |
|-----------|----------|---------|
| 333445555 | Franklin | Wong |
| 987654321 | Jennifer | Wallace |

□ ANTI JOIN: $T_1 \setminus_{X = Y} T_2$

Return the information of all the employees who *do not* work in any department which was managed since 1990.

EMPLOYEE $\setminus_{Dno=Dnumber}$ ($\sigma_{Mgr_start_date \geq '1990-01-01'}$ (**DEPARTMENT**))

```
SELECT *
FROM EMPLOYEE
WHERE Dno NOT IN ( SELECT Dnumber
    FROM DEPARTMENT
    WHERE Mgr_start_date >= '1990-01-01');
```

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|-----|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |

IN vs. NOT IN

- **IN:** comparison operator, which compares a value v with a set (or multiset) of values V and evaluates to TRUE if v is one of the elements in V .
 - V : a nested query or an explicit set
 - = ANY (or = SOME) operator: equivalent to IN operator
- **NOT IN:** TRUE if v is not any element in V .

List Essns of all employees who work the same (project, hours) combination on some project that employee whose Ssn is '123456789' works on

```
SELECT DISTINCT Essn  
FROM WORKS_ON  
WHERE (Pno, Hours) IN (SELECT Pno, Hours  
                 FROM WORKS_ON  
                 WHERE Essn = '123456789');
```

| Essn |
|-----------|
| 123456789 |

IN vs. NOT IN

List the project numbers and names of projects that have an employee with last name 'Smith' involved as worker or as manager of the department controlling the projects.

SELECT Pnumber, Pname

FROM PROJECT

WHERE Pnumber **IN** (**SELECT** DISTINCT Pno

FROM WORKS_ON

WHERE Essn **IN** (**SELECT** Ssn

FROM EMPLOYEE

WHERE Lname = 'Smith'))

OR Pnumber **IN** (**SELECT** Pnumber

FROM PROJECT, DEPARTMENT, EMPLOYEE

WHERE Dnum = Dnumber

AND Mgr_ssn = Ssn

AND Lname = 'Smith');

| Pnumber | Pname |
|---------|----------|
| 1 | ProductX |
| 2 | ProductY |

IN vs. NOT IN

- It is also possible to use an *explicit (enumerated) set of values* with IN comparison operator.

Retrieve the SSNs of all employees who work on project numbers 1, 2, or 3.

```
SELECT DISTINCT ESSN  
FROM WORKS_ON  
WHERE PNO IN (1, 2, 3);
```

```
SELECT DISTINCT ESSN  
FROM WORKS_ON  
WHERE PNO = 1 OR PNO = 2 OR PNO = 3;
```

| ESSN |
|-----------|
| 123456789 |
| 453453453 |
| 333445555 |
| 666884444 |

❑ UNION: R U S

- *Union compatibility:* R and S has the same degree and each corresponding pair of attributes has the same domain.
- Produce a new relation that includes all tuples that are either in R or in S or in both R and S. Duplicate tuples are eliminated.

| RESULT1 | SSN |
|---------|-----------|
| | 123456789 |
| | 333445555 |
| | 666884444 |
| | 453453453 |

| RESULT2 | SSN |
|---------|-----------|
| | 333445555 |
| | 888665555 |

RESULT1 U RESULT2

| RESULT | SSN |
|--------|-----------|
| | 123456789 |
| | 333445555 |
| | 666884444 |
| | 453453453 |
| | 888665555 |

**SELECT * FROM RESULT1
UNION
SELECT * FROM RESULT2;**

Check RESULT for
UNION ALL?

❑ UNION: R U S

Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT      DISTINCT PNUMBER  
  FROM       PROJECT, DEPARTMENT, EMPLOYEE  
 WHERE      DNUM=DNUMBER AND MGR_SSN=SSN  
           AND LNAME='Smith')  
  
UNION  
(SELECT      DISTINCT PNUMBER  
  FROM       PROJECT, WORKS_ON, EMPLOYEE  
 WHERE      PNUMBER=PNO AND ESSN=SSN  
           AND LNAME='Smith');
```

| PNUMBER |
|---------|
| 1 |
| 2 |

❑ INTERSECTION: $R \cap S$

USE **IN** OPERATOR ALTERNATIVELY
(DBMS-specific implementation)

- *Union compatibility*: R and S has the same degree and each corresponding pair of attributes has the same domain.
- Produce a new relation that includes all tuples that are in both R and S.

| STUDENT | FN | LN |
|---------|---------|----|
| Susan | Yao | |
| Ramesh | Shah | |
| Johnny | Kohler | |
| Barbara | Jones | |
| Amy | Ford | |
| Jimmy | Wang | |
| Ernest | Gilbert | |

STUDENT \cap INSTRUCTOR

| FN | LN |
|--------|------|
| Susan | Yao |
| Ramesh | Shah |

| INSTRUCTOR | FNAME | LNAME |
|------------|---------|-------|
| John | Smith | |
| Ricardo | Browne | |
| Susan | Yao | |
| Francis | Johnson | |
| Ramesh | Shah | |

SELECT * FROM STUDENT

INTERSECT

SELECT * FROM INSTRUCTOR;

Check RESULT for **INTERSECT ALL**?

(DBMS-specific implementation)

□ DIFFERENCE: R – S

USE **NOT IN** OPERATOR ALTERNATIVELY
(DBMS-specific implementation)

- *Union compatibility*: R and S has the same degree and each corresponding pair of attributes has the same domain.
- Produce a new relation that includes all tuples that are in R but not in S.

| STUDENT | FN | LN |
|---------|---------|---------|
| | Susan | Yao |
| | Ramesh | Shah |
| | Johnny | Kohler |
| | Barbara | Jones |
| | Amy | Ford |
| | Jimmy | Wang |
| | Ernest | Gilbert |

STUDENT - INSTRUCTOR

| FN | LN |
|---------|---------|
| Johnny | Kohler |
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |
| Ernest | Gilbert |

SELECT * FROM STUDENT

MINUS

SELECT * FROM INSTRUCTOR;

| INSTRUCTOR | FNAME | LNAME |
|------------|---------|---------|
| | John | Smith |
| | Ricardo | Browne |
| | Susan | Yao |
| | Francis | Johnson |
| | Ramesh | Shah |

INSTRUCTOR - STUDENT

| FNAME | LNAME |
|---------|---------|
| John | Smith |
| Ricardo | Browne |
| Francis | Johnson |

SELECT * FROM INSTRUCTOR

MINUS

SELECT * FROM STUDENT;

□ DIVISION: $T(Y) = R(Z) \div S(X)$, where $X \subseteq Z$

- Produce a new relation T each tuple of which appears in R in combination with every tuple in S => a sequence of π , \times , $-$

The projects of all employees

| SSN_PNOS | ESSN | PNO |
|-----------|------|-----|
| 123456789 | 1 | |
| 123456789 | 2 | |
| 666884444 | 3 | |
| 453453453 | 1 | |
| 453453453 | 2 | |
| 333445555 | 2 | |
| 333445555 | 3 | |
| 333445555 | 10 | |
| 333445555 | 20 | |
| 999887777 | 30 | |
| 999887777 | 10 | |
| 987987987 | 10 | |
| 987987987 | 30 | |
| 987654321 | 30 | |
| 987654321 | 20 | |
| 888665555 | 20 | |

Smith's projects

| SMITH_PNOS | PNO |
|------------|-----|
| | 1 |
| | 2 |

Return the employees who worked on *all* the projects that Smith worked:

$SSN_PNOS \div SMITH_PNOS$

$T1 = \pi_Y(R)$

$T2 = \pi_Y(T1 \times S - R)$

$T = T1 - T2$

| SSNS | SSN |
|------|-----------|
| | 123456789 |
| | 453453453 |

**CREATE TABLE T1 AS
SELECT DISTINCT ESSN
FROM SSN_PNOS;**

**CREATE TABLE T2 AS
SELECT DISTINCT ESSN
FROM (SELECT * FROM T1, SMITH_PNOS
MINUS
SELECT * FROM SSN_PNOS);**

**SELECT * FROM T1
MINUS
SELECT * FROM T2;**

❑ AGGREGATION: \sum *<a grouping attribute list>* \sum *<function list>* (R)

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

How many employees work in the company?

$\sum \text{COUNT Ssn}$ (EMPLOYEE)

How many employees work in each department of the company?

$\text{Dno} \sum \text{COUNT Ssn}$ (EMPLOYEE)

```
SELECT COUNT(Ssn) AS COUNT_Ssn  
FROM EMPLOYEE;
```

| COUNT_Ssn |
|-----------|
| 8 |

```
SELECT Dno, COUNT(Ssn) AS COUNT_Ssn  
FROM EMPLOYEE  
GROUP BY Dno;
```

| Dno | COUNT_Ssn |
|-----|-----------|
| 5 | 4 |
| 4 | 3 |
| 1 | 1 |

SELECT Statement – Substring Pattern Matching

- *Substring pattern matching*: allows comparison conditions on only parts of a character string, using the LIKE comparison operator.
- Partial strings are specified by using any of two reserved characters: % and _ (*underscore*)
 - %: replaces an arbitrary number of zero or more characters
 - _ : replaces a single character

```
SELECT Ssn, Fname, Lname, Address  
FROM EMPLOYEE  
WHERE Address LIKE '%Houston, TX%';
```

| Ssn | Fname | Lname | Address |
|-----------|----------|---------|--------------------------|
| 123456789 | John | Smith | 731 Fondren, Houston, TX |
| 333445555 | Franklin | Wong | 638 Voss, Houston, TX |
| 453453453 | Joyce | English | 5631 Rice, Houston, TX |
| 888665555 | James | Borg | 450 Stone, Houston, TX |
| 987987987 | Ahmad | Jabbar | 980 Dallas, Houston, TX |

```
SELECT Ssn, Fname, Lname, BDate  
FROM Employee  
WHERE BDate LIKE '_ _ _ _ _ 0 1 _ _ _';
```

| Ssn | Fname | Lname | BDate |
|-----------|--------|--------|------------|
| 123456789 | John | Smith | 1965-01-09 |
| 999887777 | Alicia | Zelaya | 1968-01-19 |

SELECT Statement – Arithmetic Operators

- Standard arithmetic operators: +, -, *, /

Show the resulting salaries if every employee working on “ProductX” project is given 10% raise.

```
SELECT FNAME, LNAME, 1.1*Salary AS INC_SAL  
FROM Employee, Works_on, Project  
WHERE SSN=ESSN  
AND PNO=PNUMBER  
AND PNAME='ProductX';
```

| FNAME | LNAME | INC_SAL |
|-------|---------|-----------|
| John | Smith | 33000.000 |
| Joyce | English | 27500.000 |

SELECT Statement - NULL

- Comparison operators for NULL values
 - IS NULL: checks whether an attribute value is NULL
 - IS NOT NULL: checks whether an attribute value is NOT NULL

Retrieve Ssn, names, Super_ssn of all employees who *do not have* supervisors.

SELECT Ssn, Fname, Lname, Super_SSN

FROM EMPLOYEE

WHERE Super_SSN IS NULL;

| Ssn | Fname | Lname | Super_SSN |
|-----------|-------|-------|-----------|
| 888665555 | James | Borg | NULL |

Retrieve Ssn, Fname, Lname, Super_SSN of all employees who *have* supervisors.

SELECT Ssn, Fname, Lname, Super_SSN

FROM EMPLOYEE

WHERE Super_SSN IS NOT NULL;

| Ssn | Fname | Lname | Super_SSN |
|-----------|----------|---------|-----------|
| 123456789 | John | Smith | 333445555 |
| 333445555 | Franklin | Wong | 888665555 |
| 453453453 | Joyce | English | 333445555 |
| 666884444 | Ramesh | Narayan | 333445555 |
| 987654321 | Jennifer | Wallace | 888665555 |
| 987987987 | Ahmad | Jabbar | 987654321 |
| 999887777 | Alicia | Zelaya | 987654321 |

SELECT Statement – More Comparison Operators

- =, >, >=, <, <=, <>, BETWEEN
 - $X \text{ BETWEEN } v1 \text{ AND } v2: X \geq v1 \text{ AND } X \leq v2$
- ANY, SOME, ALL
 - The comparison condition ($v > \text{ALL } V$) returns TRUE if the value v is greater than all the values in the set (or multiset) V .

Retrieve Ssn Fname, Lname, Dnumber, and Salary of each employee in department 5 whose salary is between \$30,000 and \$40,000.

SELECT Ssn, Fname, Lname, Dno Dnumber, Salary

FROM EMPLOYEE

WHERE Salary **BETWEEN** 30000 **AND** 40000 **AND** Dno=5;

| Ssn | Fname | Lname | Dnumber | Salary |
|-----------|----------|---------|---------|----------|
| 123456789 | John | Smith | 5 | 30000.00 |
| 333445555 | Franklin | Wong | 5 | 40000.00 |
| 666884444 | Ramesh | Narayan | 5 | 38000.00 |

SELECT Statement – More Comparison Operators

- $=, >, \geq, <, \leq, \neq, \text{BETWEEN}$
- ANY, SOME, ALL
 - The comparison condition ($v > \text{ALL } V$) returns TRUE if the value v is greater than all the values in the set (or multiset) V .

Retrieve Ssn, Fname, Lname, and Salary of each employee whose salary is greater than the salary of all employees in department 5.

SELECT Ssn, Fname, Lname, Salary

FROM EMPLOYEE

WHERE Salary **> ALL** (**SELECT** Salary **FROM** Employee **WHERE** Dno=5);

| Dno | MAX_Salary |
|-----|------------|
| 5 | 40000.00 |

SELECT Ssn, Fname, Lname, Salary

FROM EMPLOYEE

WHERE Salary **>(SELECT MAX(Salary) FROM Employee WHERE Dno=5)**;

| Ssn | Fname | Lname | Salary |
|-----------|----------|---------|----------|
| 888665555 | James | Borg | 55000.00 |
| 987654321 | Jennifer | Wallace | 43000.00 |

SELECT Statement – Multiway Join

- It is also possible to *nest join specifications*; that is, one of the tables in a join may itself be a joined table.
- Nesting allows the specification of the join of three or more tables as a single joined table, called a *multiway join*.

For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT PNUMBER, DNUM, LNAME, BDATE, ADDRESS  
FROM PROJECT, DEPARTMENT, EMPLOYEE  
WHERE DNUM=DNUMBER AND MGR_SSN=SSN AND PLOCATION='Stafford';
```

```
SELECT PNUMBER, DNUM, LNAME, BDATE, ADDRESS  
FROM (PROJECT JOIN DEPARTMENT ON DNUM=DNUMBER)  
         JOIN EMPLOYEE ON MGR_SSN=SSN  
WHERE PLOCATION='Stafford';
```

| PNUMBER | DNUM | LNAME | BDATE | ADDRESS |
|---------|------|---------|------------|-------------------------|
| 10 | 4 | Wallace | 1941-06-20 | 291 Berry, Bellaire, TX |
| 30 | 4 | Wallace | 1941-06-20 | 291 Berry, Bellaire, TX |

SELECT Statement – Nested Queries

- ❑ **Nested queries:** complete [select-from-where-...] query blocks *within* another query which is called the *outer query*.
- ❑ These nested queries can also appear in the WHERE clause or the FROM clause or the SELECT clause or other SQL clauses as needed.
- ❑ If a nested query returns a single attribute and a single tuple, the result will be a single (*scalar*) value. In such cases, it is permissible to use = instead of IN for the comparison operator.
- ❑ In general, the nested query will return a *table*, which is a set or multiset of tuples.
- ❑ SQL allows the use of tuples of values in comparisons by placing them within *parentheses*.

List the name of the department where Smith works.

SELECT Dname

FROM DEPARTMENT

WHERE Dnumber **IN** (**SELECT** Dno **FROM** EMPLOYEE **WHERE** Lname = 'Smith');

| |
|----------|
| Dname |
| Research |

SELECT Statement – Nested Queries

Retrieve Ssn, name, and address of all employees whose salary is greater than the average salary of the employees in 'Research' department.

```
SELECT SSN, FNAME, LNAME, ADDRESS  
FROM EMPLOYEE  
WHERE SALARY > (SELECT AVG(SALARY)  
                  FROM EMPLOYEE  
                  WHERE DNO IN (SELECT DNUMBER  
                           FROM DEPARTMENT  
                           WHERE DNAME='Research'));
```

| SSN | FNAME | LNAME | ADDRESS |
|-----------|----------|---------|--------------------------|
| 333445555 | Franklin | Wong | 638 Voss, Houston, TX |
| 666884444 | Ramesh | Narayan | 975 Fire Oak, Humble, TX |
| 888665555 | James | Borg | 450 Stone, Houston, TX |
| 987654321 | Jennifer | Wallace | 291 Berry, Bellaire, TX |

For each department that has more than two employees, retrieve the department number and the number of its employees who are making more than \$30,000.

```
SELECT Dno, count(*) ENumber  
FROM EMPLOYEE  
WHERE Salary>30000 AND Dno IN (SELECT Dno  
                           FROM EMPLOYEE  
                           GROUP BY Dno  
                           HAVING count(*) > 2)  
GROUP BY Dno;
```

| Dno | ENumber |
|-----|---------|
| 4 | 1 |
| 5 | 2 |

SELECT Statement – Correlated Nested Queries

- If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query*, the two queries are said to be *correlated*.
- A *correlated query* is the query where the nested query is evaluated *once for each tuple (or combination of tuples) in the outer query*.

List Ssn, name, and department number of each employee who has the salary greater than the average salary of the employees in his/her department.

SELECT Ssn, Fname, Lname, Dno Dnumber

FROM EMPLOYEE E

WHERE Salary > (**SELECT** AVG(Salary)

FROM EMPLOYEE

WHERE Dno = E.Dno);

| Ssn | Fname | Lname | Dnumber |
|-----------|----------|---------|---------|
| 333445555 | Franklin | Wong | 5 |
| 666884444 | Ramesh | Narayan | 5 |
| 987654321 | Jennifer | Wallace | 4 |

SELECT Statement – Correlated Nested Queries

Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT FNAME, LNAME  
FROM EMPLOYEE AS E  
WHERE E.SSN IN (SELECT ESSN  
         FROM DEPENDENT  
         WHERE ESSN=E.SSN  
         AND DEPENDENT_NAME=E.FNAME);
```

| FNAME | LNAME |
|-------|-------|
| | |

Retrieve Ssn and name of each employee who has two or more dependents.

```
SELECT Ssn, Lname, Fname  
FROM EMPLOYEE  
WHERE (SELECT count(*)  
         FROM DEPENDENT  
         WHERE Essn = Ssn) >= 2;
```

| Ssn | Lname | Fname |
|-----------|-------|----------|
| 123456789 | Smith | John |
| 333445555 | Wong | Franklin |

INSERT Statement

- ❑ In its simplest form, INSERT is used to add one row in a table.
- ❑ Table name and a list of values for a new row need to be specified.
- ❑ Attribute values must be listed *in the same order* as the attributes were specified in the CREATE TABLE statement.
- ❑ NOTE: All the DDL *constraints* on the table are checked for any modification on the table.
 - A *DBMS* that fully implements SQL should *support and enforce all the integrity constraints* that can be specified in the *DDL*.

INSERT Statement

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|-----|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

```
INSERT INTO EMPLOYEE  
VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30',  
'98 Oak Forest, Katy, TX', 'M', 37000, '653298653', 4);
```

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|-----|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Richard | K | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M | 37000.00 | 653298653 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

INSERT Statement

- A second form of the INSERT statement allows the user to specify explicit attribute names that correspond to the values provided in the INSERT command.
- This is useful if a relation has many attributes but only a few of those attributes are assigned values in the new tuple.
 - The values must include all attributes with NOT NULL specification and no default value.
 - Attributes with NULL allowed or DEFAULT values are the ones that can be left out.

INSERT Statement

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|-----|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Richard | K | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M | 37000.00 | 653298653 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Peter', 'Pan', 1, '123123123');
```

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|------|----------|-----------|-----|
| Peter | NULL | Pan | 123123123 | NULL | NULL | NULL | NULL | NULL | 1 |
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Richard | K | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M | 37000.00 | 653298653 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

INSERT Statement

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno)
VALUES ('Jack', 'Sparrow', 1);
```

Error Code: 1364. Field 'SSN' doesn't have a default value

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Jack', 'Sparrow', 1, '123456789');
```

Error Code: 1062. Duplicate entry '123456789' for key 'PRIMARY'

```
INSERT INTO EMPLOYEE
VALUES ('James', 'J', 'Watson', '653298653', '1962-12-30',
'98 Oak Forest, Houston, TX', 'M', 37000, '653298653', NULL);
```

Error Code: 1048. Column 'DNO' cannot be null

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Jack', 'Sparrow', 2, '234234234');
```

Error Code: 1452. Cannot add or update a child row: a foreign key constraint fails
(`company`.`employee`, CONSTRAINT `employee_ibfk_2` FOREIGN KEY (`DNO`) REFERENCES `department` (`DNUMBER`))

INSERT Statement

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno)
VALUES ('Jack', 'Sparrow', 1);
```

Error Code: 1364. Field 'SSN' doesn't have a default value

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Jack', 'Sparrow', 1, '123456789');
```

Error Code: 1062. Duplicate entry '123456789' for key 'PRIMARY'

```
INSERT INTO EMPLOYEE
VALUES ('James', 'J', 'Watson', '653298653', '1962-12-30',
'98 Oak Forest, Houston, TX', 'M', 37000, '653298653', NULL);
```

Error Code: 1048. Column 'DNO' cannot be null

```
INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
VALUES ('Jack', 'Sparrow', 2, '234234234');
```

Error Code: 1452. Cannot add or update a child row: a foreign key constraint fails

HOW TO FIX THEM?

INSERT Statement

- A variation of the INSERT command *inserts multiple tuples into a table* in conjunction with creating the table and loading it with the result of a query.

```
CREATE TABLE WORKS_ON_INFO (
    Emp_name          VARCHAR(15),
    Proj_name         VARCHAR(15),
    Hours_per_week   DECIMAL(3,1)
);
```

```
INSERT INTO WORKS_ON_INFO
SELECT E.Lname, P.Pname, W.Hours
FROM PROJECT P, WORKS_ON W, EMPLOYEE E
WHERE P.Pnumber = W.Pno AND W.Essn = E.Ssn;
```

INSERT Statement

- Another variation for loading data is to create a new table TNEW that has *the same attributes* as an existing table T, and load some of the data currently in T into TNEW. The syntax for doing this uses the *LIKE* clause.

```
CREATE TABLE D5EMPS LIKE EMPLOYEE;  
  
INSERT INTO D5EMPS  
SELECT *  
FROM EMPLOYEE  
WHERE Dno = 5;
```

? COMPARED with:

```
CREATE TABLE D5EMPS AS  
SELECT *  
FROM EMPLOYEE  
WHERE Dno = 5;
```

DELETE Statement

- The DELETE command removes tuples from a relation.
- It includes a WHERE clause, similar to that used in an SQL query, to select the tuples to be deleted.
- Tuples are explicitly deleted from only one table at a time. However, the deletion may propagate to tuples in other relations if referential triggered actions are specified in the referential integrity constraints of the DDL.
- A missing WHERE clause specifies that all tuples in the relation are to be deleted; however, the table remains in the database as an empty table.
 - The DROP TABLE command is used to remove the table definition.

DELETE Statement

Remove from EMPLOYEE table the details of the employees whose last name is Brown.

```
DELETE FROM EMPLOYEE  
WHERE LNAME='Brown';
```

How many rows are deleted?
Is any constraint violated?

Remove from EMPLOYEE table the details of the employee whose SSN is 123456789.

```
DELETE FROM EMPLOYEE  
WHERE SSN='123456789';
```

Remove from EMPLOYEE table the details of all the employees who work in Research department.

```
DELETE FROM EMPLOYEE  
WHERE DNO IN (SELECT DNUMBER  
                 FROM DEPARTMENT  
                 WHERE DNAME = 'Research');
```

Remove from EMPLOYEE table the details of all the employees.

```
DELETE FROM EMPLOYEE;
```

UPDATE Statement

- The UPDATE command is used to *modify attribute values of one or more selected tuples in a table.*
- As in the DELETE command, a *WHERE clause* in the UPDATE command selects the tuples to be modified from a single table.
 - Updating a primary key value may propagate to the foreign key values of tuples in other relations if such a referential triggered action is specified in the referential integrity constraints of the DDL.
 - To modify multiple relations, several UPDATE commands must be issued.
- The *SET clause* in the UPDATE command specifies the attributes to be modified and their new values.

UPDATE Statement

Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
UPDATE PROJECT  
SET PLOCATION = 'Bellaire', DNUM = 5  
WHERE PNUMBER=10;
```

Increase 2 hours for all the employees of department 1 working on the projects.

```
UPDATE WORKS_ON  
SET HOURS = HOURS + 2  
WHERE ESSN IN (SELECT DISTINCT SSN  
                 FROM EMPLOYEE  
                 WHERE DNO = 1);
```

Data Control Language - Access Control

- ❑ **Database access control:** restricting data access and database activities, e.g. restrict users from querying specified tables or executing specified database statements.
- ❑ **Privilege:** the right to perform some action, e.g. run specific SQL statements.
- ❑ **User:** an account through which you can log in to the database, and to establish the means by which *Oracle Database (DBMS)* permits access by the user
- ❑ **Role:** a *named* set of privileges that can be granted to users or to other roles

Data Control Language - Access Control

❑ Privilege categories

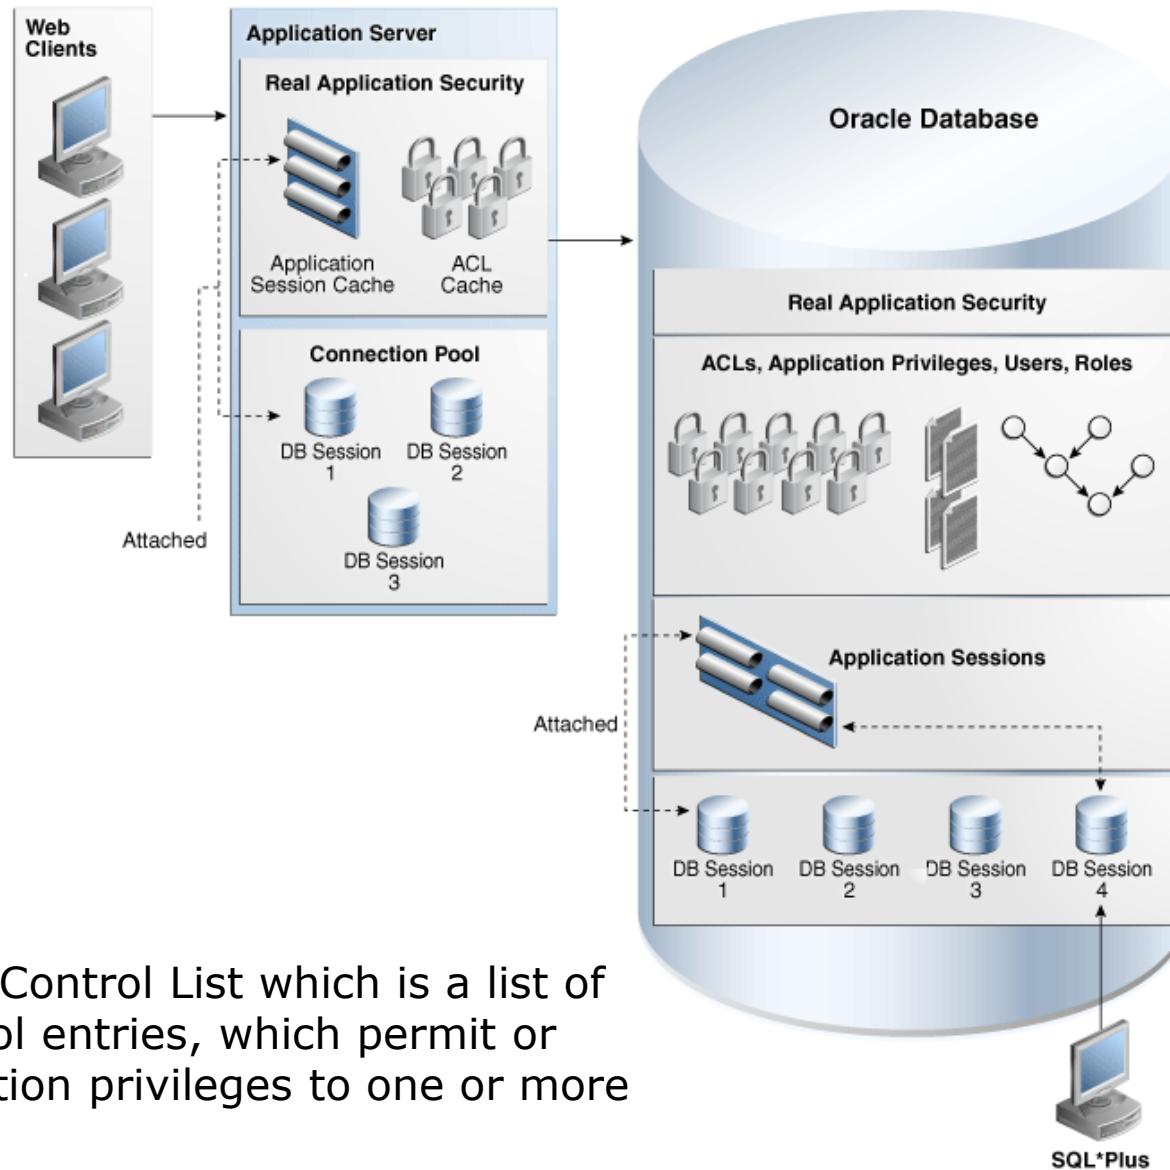
■ System privilege

- ❑ This is the right to perform a specific action in the database, or perform an action on any objects of a specific type. For example, CREATE USER and CREATE SESSION are system privileges.

■ Object privilege

- ❑ This is the right to perform a specified action on an object, for example, query the employees table. Privilege types are defined by the database.

Oracle Database Real Application Security Components – Oracle 19c



ACL: Access Control List which is a list of access control entries, which permit or deny application privileges to one or more principals.

GRANT Statement

- The GRANT statement provides privileges for a user to perform specific actions.

Create user **psmith** as a new user while granting **psmith** the **CREATE SESSION** system privilege.

If a password is specified using the **IDENTIFIED BY** clause, and the user name does not exist in the database, then a new user with that user name and password is created.

GRANT CREATE SESSION TO psmith IDENTIFIED BY *password*;

Grant the **UPDATE**, **INSERT**, and **DELETE** object privileges for all columns of the **EMPLOYEE** table to the user **psmith**.

GRANT UPDATE, INSERT, DELETE ON EMPLOYEE TO psmith;

The **WITH GRANT OPTION** clause with the **GRANT** statement can enable a grantee, **psmith**, to grant object privileges to *other* users.

GRANT SELECT ON EMPLOYEE TO psmith WITH GRANT OPTION;

REVOKE Statement

- ❑ The REVOKE SQL statement revokes system privileges and roles.
- ❑ Any user with the ADMIN option for a system privilege or role can revoke the privilege or role from any other database user or role.
- ❑ The revoker does not have to be the user that originally granted the privilege or role. Users with GRANT ANY ROLE can revoke *any* role.
 - Can only revoke the privileges that the person who granted the privilege, directly authorized.
 - Cannot revoke grants that were made by other users to whom was granted the GRANT OPTION.
 - However, there is a cascading effect. If the object privileges of the user who granted the privilege are revoked, then the object privilege grants that were propagated using the GRANT OPTION are revoked.

Revoke **SELECT** and **INSERT** privileges on the **EMPLOYEE** table from user **psmith**.
REVOKE SELECT, INSERT ON EMPLOYEE FROM psmith;

Revoke all object privileges for **DEPARTMENT** table that was originally granted to the **human_resources** role.

REVOKE ALL ON DEPARTMENT FROM human_resources;

GRANT and REVOKE Statements

```
-- from user 'root'@'localhost'
```

```
CREATE USER 'chau2'@'localhost' IDENTIFIED BY 'chau2';
```

```
GRANT ALL PRIVILEGES ON company.* TO 'chau2'@'localhost';
```

```
GRANT CREATE, ALTER, SELECT ON test.* TO 'chau2'@'localhost';
```

```
SELECT * FROM 'demo_20201003'.'employee';
```

Error Code: 1142. SELECT command denied to user

'chau2'@'localhost' for table 'employee'

```
GRANT SELECT ON demo_20201003.* TO 'chau2'@'localhost';
```

```
SHOW GRANTS FOR 'chau2'@'localhost';
```

```
REVOKE ALL PRIVILEGES ON *.* FROM 'chau2'@'localhost';
```

Data Control Language – Transaction Control

- ▣ **Transaction:** a logical data processing unit on a database
 - Change a database from a consistent state to another consistent one
 - ACID properties: atomic, consistent, isolated, durable
- ▣ **Transaction Control Language:** a set of commands which are used to mark the synchronization points (*sync points*) of transactions
 - COMMIT
 - ROLLBACK
 - SET AUTOCOMMIT OFF

Data Control Language – Transaction Control

- ❑ **COMMIT**: confirms the operations in the transaction have successfully completed. It is the responsibility of the DBMS to make sure that all those confirmed operations are permanent.
- ❑ **ROLLBACK**: cancels all the operations from the previous sync point to the ROLLBACK point.
- ❑ **SET AUTOCOMMIT OFF**: equivalent to “begin transaction”, opposite to SET AUTOCOMMIT ON. When AUTOCOMMIT is ON, each SQL command is considered as a separate transaction and is committed automatically after its execution. To group several SQL commands into one transaction, AUTOCOMMIT must be OFF.

Database Stored Procedures and SQL/PSM

- Database program modules—procedures or functions—that are *stored and executed by the DBMS at the database server.*
 - These are historically known as *database stored procedures*, although they can be **functions** or **procedures**.
 - The term used in the SQL standard for stored procedures is *persistent stored modules* because these programs are stored persistently by the DBMS, similarly to the persistent data stored by the DBMS.

Database Stored Procedures and SQL/PSM

- Database stored procedures are useful in the following circumstances:
 - If a database program is needed by *several applications*, it can be stored at the server and invoked by any of the application programs. This reduces duplication of effort and improves software modularity.
 - Executing a program at the server can *reduce data transfer and communication cost* between the client and server in certain situations.
 - These procedures can enhance the modeling power provided by views by allowing *more complex types of derived data* to be made available to the database users via the stored procedures. Additionally, they can be used to check for *complex constraints* that are beyond the specification power of assertions and triggers.

Database Stored Procedures and SQL/PSM

- CREATE, ALTER, and DROP statements create, modify, and drop a stored procedure or a stored function.
- Functions/procedures can be written in SQL itself, or in an external programming language (e.g., C, Java).
 - Loops, if-then-else, assignment
- Some database systems (e.g., Oracle, MS SQL Server) support **table-valued functions**, which can return a table as a result.
- Functions are **parameterized views** that generalize the regular notion of views by allowing *parameters*.

Stored Functions

- Given a department name, return the number of employees who work for the corresponding department.

```
create function dept_count (dept_name varchar (15))
returns int
begin
    declare d_count int;
    select count(*) into d_count
    from employee join department on dno = dnumber
    where dname = dept_name;
    return d_count;
end
```

```
select *
from department
where dept_count (dname) > 10;
```

| DNAME | DNUMBER | MGR_SSN | MGR_START_DATE |
|----------------|---------|-----------|----------------|
| Administration | 4 | 987654321 | 1995-01-01 |
| Research | 5 | 333445555 | 1988-05-22 |

Stored Functions

```
CREATE FUNCTION [dbo].[dept_count]
(
    @dept_name varchar(15)
)
RETURNS int
AS
BEGIN

    DECLARE @d_count int;

    SET @d_count = (SELECT COUNT(*)
                    FROM dbo.employee join dbo.department on dno = dnumber
                    WHERE dname = @dept_name);

    -- Return the result of the function
    RETURN @d_count;

END
```

Stored Functions

- Return the department number, department name, and average salary of each department of more than n employees.

```
create function dept_avg_salary (n int)
returns d_a_s TABLE (
    dname          varchar (15),
    dnumber        int,
    avg_sal        decimal (10,2)
)
begin
    insert into d_a_s
    select dname, dnumber, AVG(salary)
    from employee join department on dno = dnumber
    where dept_count(dname) > n
    group by dnumber, dname;
    return;
end
```

```
select *
from dept_avg_salary (2);
```

| dname | dnumber | avg_sal |
|----------------|---------|----------|
| Administration | 4 | 31000.00 |
| Research | 5 | 33250.00 |

Stored Functions

```
CREATE FUNCTION [dbo].[dept_avg_salary]
(
    @empnumber int
)
RETURNS
@d_a_s TABLE
(
    dname varchar(15),
    dnumber int,
    avg_sal decimal(10, 2)
)
AS
BEGIN
    insert into @d_a_s
    select dname, dnumber, AVG(salary)
    from dbo.employee join dbo.department on dno = dnumber
    where dbo.dept_count(dname) > @empnumber
    group by dnumber, dname;
    RETURN
END
```

Stored Procedures

- Given a department name, return the number of employees who work for the corresponding department.

```
CREATE PROCEDURE [dbo].[p_dept_count]
    (@dept_name varchar (15), @d_count int OUT)
AS
BEGIN
    SET @d_count = (select COUNT(*)
                    from employee join department on dno = dnumber
                    where dname = @dept_name);
END
```

```
declare @d_count int;
exec dbo.p_dept_count 'Research', @d_count output;
print @d_count;
```

4

Stored Procedures

- Return the department number, department name, and average salary of each department of more than n employees.

```
CREATE PROCEDURE [dbo].[p_dept_avg_salary] (@n INT)
AS
BEGIN
    select dname, dnumber, AVG(salary)
    from dbo.employee join dbo.department on dno = dnumber
    where dbo.dept_count(dname) > @n
    group by dnumber, dname;
END
```

```
DECLARE @dept_avg_sal TABLE (dname varchar(15), dnumber int, avg_salary decimal(10,2));
INSERT INTO @dept_avg_sal EXEC dbo.p_dept_avg_salary 2;
SELECT * FROM @dept_avg_sal;
```

| dname | dnumber | avg_sal |
|----------------|---------|----------|
| Administration | 4 | 31000.00 |
| Research | 5 | 33250.00 |

Procedural language constructs

- ❑ SQL supports constructs that gives it almost all the power of a general-purpose programming language.
 - **Warning:** most database management systems implement their own variant of the standard syntax.
- ❑ Compound statement: **begin ... end**,
 - Multiple SQL statements are given between **begin** and **end**.
 - Local variables can be declared within a compound statements.
- ❑ **for, while** and **repeat** statements
- ❑ **if, if ... else ...** statements
- ❑ **case** statements, ...

Procedural language constructs

```
[begin_label:] BEGIN  
  [statement_list]  
END [end_label]
```

```
DECLARE var_name [, var_name] ... type [DEFAULT value]
```

```
DECLARE cursor_name CURSOR FOR select_statement
```

```
OPEN cursor_name
```

```
FETCH cursor_name INTO var_name [, var_name] ...
```

```
CLOSE cursor_name
```

```
SET var_name = expr [, var_name = expr] ...
```

```
SELECT col_name [, col_name] ...  
  INTO var_name [, var_name] ...  
  table_expr
```

Procedural language constructs

```
IF search_condition THEN statement_list  
    [ELSEIF search_condition THEN _statement_list] ...  
    [ELSE statement_list]  
END IF
```

```
CASE case_value  
    WHEN when_value THEN statement_list  
        [WHEN when_value THEN statement_list] ...  
    [ELSE statement_list]  
END CASE
```

```
[begin_label:] LOOP  
    statement_list  
END LOOP [end_label]
```

```
[begin_label:] WHILE search_condition DO  
    statement_list  
END WHILE [end_label]
```

```
[begin_label:] REPEAT  
    statement_list  
UNTIL search_condition  
END REPEAT [end_label]
```

```
LEAVE label
```

```
ITERATE label
```

```
RETURN expr
```

Stored Procedures

- Insert data of a new employee using 20000 if salary is <=0.

```
CREATE DEFINER=`root`@`localhost` PROCEDURE `insert_employee`
  (id varchar(9), fname varchar(15), lname varchar(15), dno int(11), salary decimal(10,2))
BEGIN
  if salary <= 0.0 then
    set salary = 20000;
  end if;

  insert into employee (ssn, fname, lname, salary, dno)
  values (id, fname, lname, salary, dno);
END
```

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|------|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| fn_a | NULL | ln_b | 135792468 | NULL | NULL | NULL | 20000.00 | NULL | 1 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Richard | K | Marini | 653298653 | 1962-12-30 | 98 Oak Forest, Katy, TX | M | 37000.00 | 653298653 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

```
call insert_employee ('135792468', 'fn_a', 'ln_b', 1, 0);
```

Stored Procedures

- Insert data of a new employee using the average salary of the department if no salary is specified (<=0).

```
CREATE PROCEDURE [dbo].[insert_employee]
    (@id varchar(9), @fname varchar(15), @lname varchar(15), @dno int, @salary decimal(10, 2))
AS
BEGIN
    declare @avg_sal decimal(10, 2)

    if (@salary<=0)
        SET @avg_sal = (select avg(salary)
                          from employee
                          where dno = @dno)
    else
        SET @avg_sal = @salary

    insert into employee (ssn, fname, lname, salary, dno)
    values (@id, @fname, @lname, @avg_sal, @dno)

END
```

| FNAME | MINIT | LNAME | SSN | BDATE | ADDRESS | SEX | SALARY | SUPER_SSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|------|----------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000.00 | 333445555 | 5 |
| fn_a | NULL | In_b | 135792468 | NULL | NULL | NULL | 31000.00 | NULL | 4 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000.00 | 888665555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000.00 | 333445555 | 5 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000.00 | 333445555 | 5 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000.00 | NULL | 1 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000.00 | 888665555 | 4 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000.00 | 987654321 | 4 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000.00 | 987654321 | 4 |

```
EXEC dbo.insert_employee '135792468', 'fn_a', 'In_b', 4, 0;
```

Triggers

- A trigger is a named database object that is associated with a table, and that *automatically* activates when a particular event occurs for the table.
 - Activation time
 - BEFORE or AFTER an event
 - *INSTEAD OF* an event
 - Event: INSERT, DELETE, UPDATE
 - Refer to columns in the table associated with the trigger by using the aliases OLD (**deleted**) and NEW (**inserted**).
 - OLD.col_name: a column of an existing row before it is updated or deleted
 - NEW.col_name: a column of a new row to be inserted or an existing row after it is updated

Triggers

□ When to use triggers

- Enforce application-based semantic constraints
- Maintain summary data
- Replicate databases by recording changes to other special tables

□ When not to use triggers **Cascading execution**

- Encapsulation facilities available for update (insert, delete, update) methods with constraints
- Materialized views available for maintaining summary data
- Built-in DBMS support ready to replicate databases

Triggers

```
CREATE  
[DEFINER = { user | CURRENT_USER }]  
TRIGGER trigger_name trigger_time trigger_event  
ON tbl_name FOR EACH ROW trigger_body
```

Create a trigger to ensure each employee works on at least one project.

DELIMITER \$\$

```
CREATE TRIGGER trgTotalEmployee_after_delete_works_on AFTER DELETE ON works_on  
FOR EACH ROW  
BEGIN
```

```
declare total_participation condition for SQLSTATE '45000';
```

```
IF NOT EXISTS (select * from works_on where essn = old.essn) THEN  
    SIGNAL total_participation  
    SET MESSAGE_TEXT = 'Every employee must work on at least one project!';  
END IF;
```

```
END $$
```

Event 34 of type Error at 11:08:11

Action: delete from works_on where essn = '888665555' and pno = 20

Message: Error Code: 1644. Every employee must work on at least one project!

Triggers

- Update the number of employees of the department from which an employee is removed.

```
CREATE TRIGGER delete_trigger ON employee AFTER DELETE
AS
BEGIN
    DECLARE @dname VARCHAR(15);
    DECLARE @count INT;

    SELECT @dname = dname, @count = COUNT(*)
    FROM department join employee ON dnumber = dno
    WHERE dno IN (SELECT distinct dno FROM deleted)
    GROUP BY dname;

    PRINT 'Department name: ' + @dname + ' -- Number of employees: ' + CAST(@count AS VARCHAR(15));
END
```

Triggers

- Check if a new employee is with the same department as his/her supervisor before inserted.

```
CREATE TRIGGER insert_emp_trigger ON employee INSTEAD OF INSERT
AS
BEGIN
    IF (EXISTS(SELECT * FROM employee WHERE ssn = (select super_ssn from inserted)
                AND dno = (select dno from inserted)))
        BEGIN
            INSERT INTO employee
            SELECT * FROM inserted;
        END
    END
```

NOTES on Functions, Procedures, and Triggers

- More details of stored functions, stored procedures, and triggers
 - Check the specification of a particular DBMS
 - See more examples on MS SQL Server with a demonstration uploaded in BKEL

Summary

❑ SQL = Structured Query Language

- Data Definition Language
- Data Manipulation Language
- Data Control Language
- Stored Functions
- Stored Procedures
- Triggers

ANSI/ISO Standard:
SQL-86,
SQL-89,
SQL-92,
SQL-1999,
SQL-2003,
SQL-2006,
SQL-2008,
SQL-2011,
SQL-2016,
SQL-2019

- A *declarative* non-procedural language, based on tuple relational calculus, supported by existing RDBMSs
- Associated with relational algebra for *internal* representation, processing, and optimization

Summary

❑ Further readings

- DDL statements for views, indexes
- Full syntax of each SQL statement on a specific relational DBMS
 - ❑ Oracle
 - ❑ MS SQL Server
 - ❑ MySQL
 - ❑ PostgreSQL
 - ❑ SQLite
 - ❑ ...

Chapter 4: The SQL Language



Review

- 1. Given the Company database, write a SELECT statement for each requirement:
 - 1.1. Retrieve Ssn, name, and address of the employees who work in department 5.
 - 1.2. Retrieve Ssn, name, and address of the employees who work in 'Research' department.
 - 1.3. Retrieve Ssn, name, and address of the employees who work in departments in Houston.
 - 1.4. Return the average salary of the employees who work in department 1.
 - 1.5. Return the maximum salary of the employees in each department.
 - 1.6. Return the number of the employees who have salaries greater than the average salary of the employees who work in department 1.

Review

- 1. Given the Company database, write a SELECT statement for each requirement:
 - 1.7. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.
 - 1.8. Find the names of employees who work on all the projects controlled by department number 5.
 - 1.9. Make a list of project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.
 - 1.10. List the names of all employees with two or more dependents in department 5.
 - 1.11. Retrieve the names of employees who have no dependents.

Review

- 1. Given the Company database, write a SELECT statement for each requirement:
 - 1.12. List the names of managers who have at least one dependent.
 - 1.13. Retrieve the names of all employees in department 5 who work more than 10 hours per week on the ProductX project.
 - 1.14. List the names of all employees who have a dependent with the same first name as themselves.
 - 1.15. Find the names of all employees who are directly supervised by 'Franklin Wong'.
 - 1.16. For each project, list the project name and the total hours per week (by all employees) spent on that project.
 - 1.17. Retrieve the names of all employees who work on every project.

Review

- 1. Given the Company database, write a SELECT statement for each requirement:
 - 1.18. Retrieve the names of all employees who do not work on any project.
 - 1.19. For each department, retrieve the department name and the average salary of all employees working in that department.
 - 1.20. Retrieve the average salary of all female employees.
 - 1.21. Find the names and addresses of all employees who work on at least one project located in Houston but whose department has no location in Houston.
 - 1.22. List the last names of all department managers who have no dependents.
 - 1.23. List the department name, the number of projects, and the number of employees of each department that controls those projects in the same location.

Review

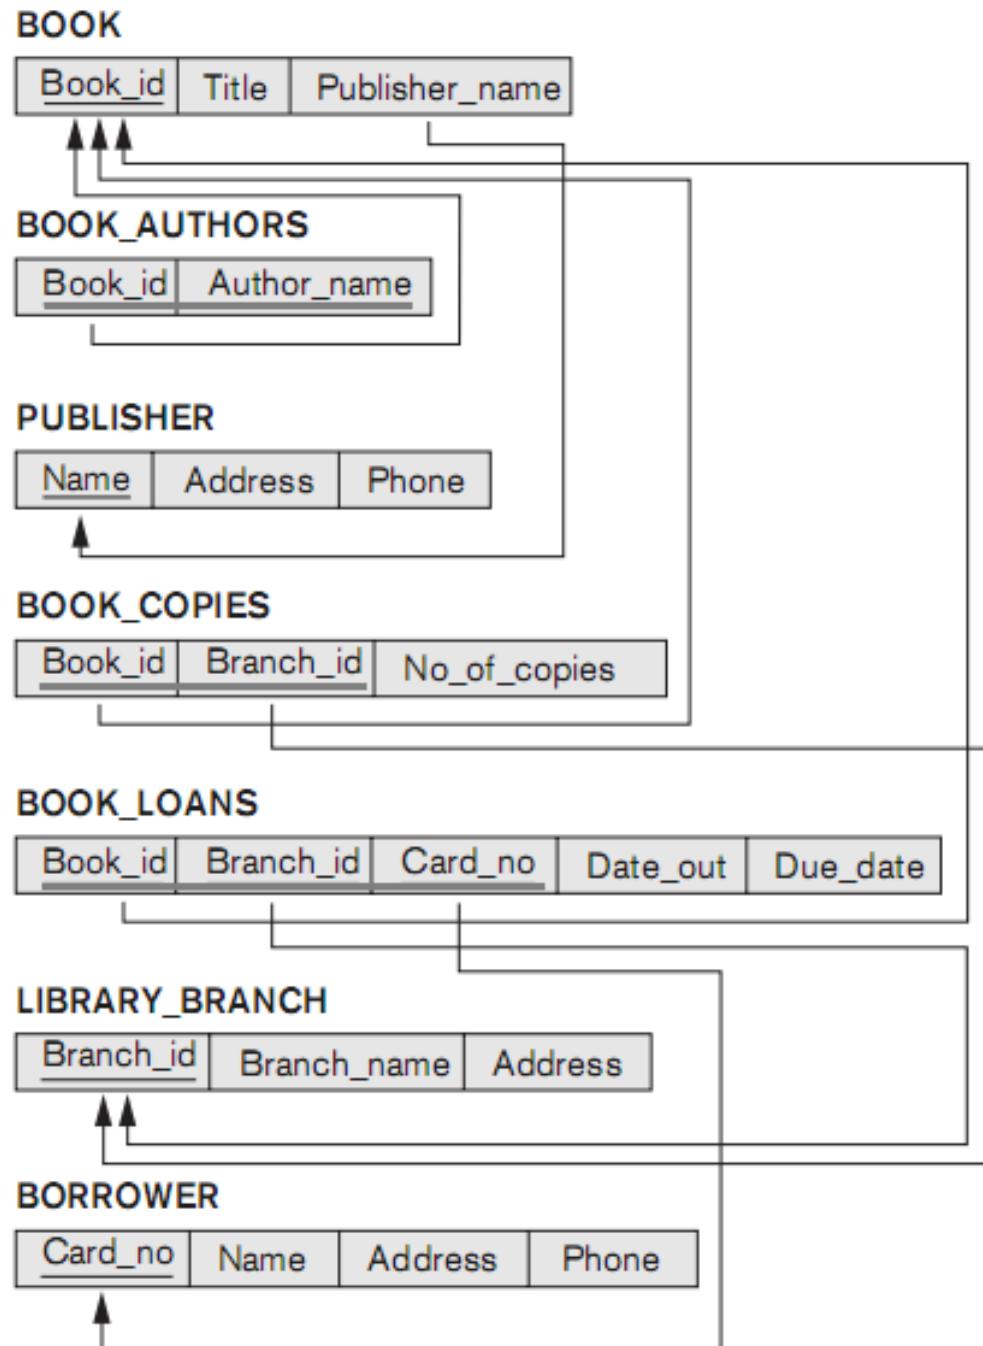
- 2. For the Company database, write a DML statement for each requirement. Is any constraint violated?
 - 2.1. Insert <'Robert', 'F', 'Scott', '943775543', '1972-06-21', '2365 Newcastle Rd, Bellaire, TX', M, 58000, '888665555', 1> into EMPLOYEE.
 - 2.2. Insert <'ProductA', 4, 'Bellaire', 2> into PROJECT.
 - 2.3. Insert <'Production', 4, '943775543', '2007-10-01'> into DEPARTMENT.
 - 2.4. Insert <'677678989', NULL, '40.0'> into WORKS_ON.
 - 2.5. Insert <'453453453', 'John', 'M', '1990-12-12', 'spouse'> into DEPENDENT.
 - 2.6. Delete the WORKS_ON tuples with Essn = '333445555'.
 - 2.7. Delete the EMPLOYEE tuple with Ssn = '987654321'.
 - 2.8. Delete the PROJECT tuple with Pname = 'ProductX'.
 - 2.9. Delete the DEPARTMENT tuple with Dname = 'Administration'.

Review

- 2. For the Company database, write a DML statement for each requirement. Is any constraint violated?
 - 2.10. Modify the Mgr_ssn and Mgr_start_date of the DEPARTMENT tuple with Dnumber = 5 to '123456789' and '2007-10-01', respectively.
 - 2.11. Modify the Super_ssn attribute of the EMPLOYEE tuple with Ssn = '999887777' to '943775543'.
 - 2.12. Modify the Hours attribute of the WORKS_ON tuple with Essn = '999887777' and Pno = 10 to '5.0'.
 - 2.13. Modify the Dname attribute of the DEPARTMENT tuple with Dname = 'Administration' to 'Headquarters'.
 - 2.14. Modify the Salary attribute of the EMPLOYEE tuples with the average salary of the department 'Research' where the salary of its manager is 10% higher.
 - 2.15. Modify the Dno attribute of the EMPLOYEE tuples with Dno = 1 to Dno = 2.

Review

- 3. Given the relational schema of the LIBRARY database, write DDL statements to implement this schema with any assumption needed for data types and constraints.



Review

- 4. Write SQL statements on the LIBRARY database in question 3.
 - 4.1. Add an ISBN attribute of varchar(13) into BOOK relation schema where ISBN is a secondary key.
 - 4.2. How many copies of the book titled The Lost Tribe are owned by the library branch whose name is 'Sharpstown'?
 - 4.3. How many copies of the book titled The Lost Tribe are owned by each library branch?
 - 4.4. Retrieve the names of all borrowers who do not have any books checked out.
 - 4.5. For each book that is loaned out from the Sharpstown branch and whose Due_date is today, retrieve the book title, the borrower's name, and the borrower's address.
 - 4.6. For each library branch, retrieve the branch name and the total number of books loaned out from that branch.
 - 4.7. Retrieve the names, addresses, and number of books checked out for all borrowers who have more than five books checked out.
 - 4.8. For each book authored (or coauthored) by Stephen King, retrieve the title and the number of copies owned by the library branch whose name is Central.

Review

- 5. Write stored functions to process the data in the LIBRARY database as follows:
 - 5.1. Return ISBN and book title of each book available in every library branch.
 - 5.2. Given a publisher name and a date, return the number of books published by this publisher borrowed on the date.
 - 5.3. Given an author name, return the details of the borrowers who borrowed his/her books the most.

Review

- 6. Write stored procedures to process the data in the LIBRARY database as follows:
 - 6.1. Given a date, update a temporary table with the details of all the books that are not available on that date because all of them were borrowed.
 - 6.2. Given an author name, return the most recent date when his/her book was borrowed.
 - 6.3. Return the library branches where the books have been borrowed the most.

Review

- 7. Write triggers to process the data in the LIBRARY database as follows:
 - 7.1. Enforce the constraint which is: the number of copies of each book in any library branch must be at least 2.
 - 7.2. Enforce the constraint which is: a borrower cannot borrow more than 3 books at the same time.
 - 7.3. Update a temporary table of the books whose copies are not available in any library branch because all of them were borrowed.

Chapter 5: Relational Database Design

- 5.1. Design Guidelines for Relation Schemas
- 5.2. Functional Dependencies
- 5.3. Normal Forms Based on Primary Keys
- 5.4. Boyce-Codd Normal Form
- 5.5. Properties of Relational Decompositions
- 5.6. Algorithms for Relational Database Schema Design